

# Contents

|                                                             |    |
|-------------------------------------------------------------|----|
| Coverage and Max-Min Diversity in Synthetic Data Generation | 1  |
| Abstract                                                    | 1  |
| Part I — Max-min diversity: no two items alike              | 1  |
| Part II — Coverage: reach as much as possible               | 23 |
| Part III — Where the two objectives disagree                | 38 |
| III.1 The dissociation, measured                            | 38 |
| III.2 Why the scoring rule has to fork                      | 38 |
| III.2b The duality, taken seriously                         | 39 |
| III.2c The comparison that proves the fork                  | 39 |
| III.3 What both objectives share                            | 40 |
| III.4 Practical guidance                                    | 40 |

## Coverage and Max-Min Diversity in Synthetic Data Generation

Two objectives, one budget, and the conditional dimension that limits both

---

### Abstract

Synthetic corpora are generated under a fixed budget of  $n$  model calls, and two different things are wanted from that budget. **Coverage** asks to reach as much of the space as possible in  $n$  turns. **Max-min** asks that no two of the  $n$  items resemble each other. These are not two phrasings of one goal: coverage will place two items near each other if between them they reach a large region, and max-min will leave most of the space empty provided nothing collides. We study both, on one generator, one set of embedders, and one budget, and we show they disagree sharply enough that a method optimizing either can be near-worst at the other — greedy k-center wins min-gap in both our domains and finishes **last** on coverage, at a sixth of random.

Our central claim is that neither objective is limited by its own optimizer. Both are limited by the **support**. Conditioned on a fixed prompt, a language model's output concentrates on a submanifold of dimension  $m$  far below the dimension  $d$  of the space it could reach, and fifteen numerical checks establish the consequences: novelty at fixed prompt decays as  $n^{-1/m}$ , not  $n^{-1/d}$ ; one prompt  $\varepsilon$ -covers a vanishing  $\varepsilon^{d-m}$  fraction; only prompt motion *transverse* to the already-occupied span raises the ceiling; and unconstrained max-min selection is inconsistent, selecting an off-manifold candidate essentially whenever one is offered.

We assume an embedding oracle and, critically, **no inverse**. We can compute exactly where the next item ought to land and have no way to decode that point into text. Every architectural choice follows: the system must propose, measure and select rather than solve, and its only steering handles are language-valued. We therefore elicit latent axes from the generator itself and rank them with an **axis-scoring rule** — spread  $\times$  transversality  $\times$  independence  $\times$  headroom — which takes two forms, one per objective, differing on exactly the two terms the objectives differ on. The scoring is evaluated recursively, in the space the artifact actually occupies, and the latent lattice refines itself where it saturates.

We validate on ~46,000 real generations from openai/gpt-5.6-luna, ~700 rendered images, and ~200 rendered instrumentals, measuring diversity at three levels: literal ( $n$ -gram), latent (text embedding), and — where the artifact is not text — in the space the product occupies (CLIP for images, CLAP and MERT for audio). Headline findings: a psychometric corpus that is **72.3% exact duplicates** with one item repeated 1,637 times, which temperature barely dents (61.6%) and latent conditioning nearly eliminates (0.2%); literal and latent diversity moving in **opposite directions** as  $n$  grows; text-embedding similarity predicting rendered-image similarity at only  $r = 0.155$ ; and, against five published methods, wins of 36% (max-min, images) and, on coverage, first place among twelve corpora against the released Alpaca, PersonaHub and

WizardLM sets — 0.4441 against Alpaca's 0.3722 at matched evaluated- $n$  on a leak-free human-written reference, at one-twentieth of Alpaca's generation budget — via retrieval-aided, density-adaptive conditioning that keeps every render. We also report what did not work, including a cross-corpus coverage metric that inverts, and an arm showing that Part I's own orthogonalization machinery, applied to the coverage objective, scores below doing nothing.

---

## Part I — Max-min diversity: no two items alike

### Abstract

We study the problem of generating a corpus that stays diverse as it grows without bound. The objective is a convex combination of two terms evaluated against everything generated so far: the mean distance from a new item's embedding to the existing corpus, which we want small because it keeps items typical and on-manifold, and the minimum distance to the existing corpus, which we want large because it keeps items non-redundant. We call the resulting problem *infinite-horizon diversity* (IHD), and we are interested specifically in its asymptotics: what happens to the objective, and to the cost of satisfying it, as  $n$  grows.

Our central claim locates the interesting structure in the **support** rather than in the objective. Conditioned on a fixed prompt  $x$ , a language model's output distribution concentrates on a submanifold of dimension  $m$  that is far smaller than the dimension  $d$  of the reachable semantic manifold. We prove and numerically verify five consequences of that gap. Novelty at fixed prompt decays like  $n^{-1/m}$ , not  $n^{-1/d}$ , so saturation arrives much sooner than the ambient dimension suggests (measured slopes  $-0.472$ ,  $-0.309$ ,  $-0.195$  for  $m = 2, 3, 5$  against predictions  $-0.5$ ,  $-0.333$ ,  $-0.2$ ). A single prompt  $\epsilon$ -covers a vanishing  $\epsilon^{d-m}$  fraction of the manifold. Only prompt motion *transverse* to the already-occupied span raises the reachable dimension: parallel motion leaves the union's effective dimension at 1.9, transverse motion takes it to 11.2. Under free prompt-switching the optimal number of samples per prompt is exactly one; an interior optimum appears only once switching has overhead ( $n^* = 10$  at cost  $c = 3$ ,  $n^* = 30$  at  $c = 30$ ). And unconstrained max-min selection is not consistent: it picks an off-manifold candidate essentially whenever one appears in the pool, an  $8.4\times$  amplification over the base rate.

We assume an embedding oracle mapping text to a fixed-size vector, and — critically — **no reverse oracle**. We can compute exactly where we would like the next item to land and still have no way to decode that point into text. Every architectural choice follows from that asymmetry: the system must propose, measure, and select rather than solve, and the only steering handles are language-valued.

We therefore build the generation stack around latent behaviors elicited from the generator itself, organized by an **axis-scoring rule** that ranks candidate latent variables by the product of four measurable factors — spread, transversality, independence, and headroom — and selects each variable's most-different values by farthest-point search in level space. When the score reports that no available variable clears a promise floor, it emits a *refine* decision, and the generator is asked to split the exhausted variable into finer sub-variables. The recursion is the only mechanism in the paper that changes the asymptote rather than the constant.

We test this on **43,000 real generations** from `openai/gpt-5.6-luna` across three domains — DALL-E instructions for post-modern artworks, psychometric exam items, and instrumental-music prompts — measuring diversity at three levels: literal ( $n$ -gram), latent (text embedding), and, where the artifact is not text, in the space the product actually occupies (CLIP for rendered images, CLAP and MERT for rendered audio). Four results stand out. Asked the same reasonable question ten thousand times, the model returns a psychometric corpus that is **72.3% exact duplicates**, one item repeated 1,637 times; raising the temperature to 1.6 leaves it at 61.6%, while latent conditioning drops it to 0.2%. Literal and latent diversity move in **opposite directions** as the corpus grows. Text-embedding similarity predicts *rendered-image* similarity at only  $r = 0.155$ , so every text-side diversity method — ours included — optimizes a proxy explaining  $\sim 2\%$  of the variance in what the user receives. And against five real published methods (Self-Instruct, Evol-Instruct, Persona-Hub, high-temperature, naive) our method wins every metric, beating the strongest by 36% in mean-centered

Vendi; against **human-written MMLU** it wins on lexical diversity and self-repetition while reaching 69% of a human exam bank's semantic diversity.

In simulation, a fixed prompt in a world with  $d = 24$  and  $m = 3$  reaches a Vendi Score of 1.06 where the full stack reaches 11.23, and recursive refinement lifts the reachable dimension from 13 to 24.

## I.1 Introduction

Ask a language model for a poem ten times and you get ten poems. Ask it ten thousand times and you do not get ten thousand poems; you get a few hundred poems and a great deal of paraphrase. Any fixed conditional distribution behaves this way under repeated sampling, whatever model supplies it. The distribution has a shape, and repeated sampling traces that shape ever more densely rather than expanding it.

The practical version of this problem is now everywhere. Synthetic training data is valuable in proportion to how much it is *not* already in the training set. Evaluation suites that cluster in a few families give false assurance. Test-item banks, red-team prompt sets, persona corpora, and augmentation pipelines all consume generated text in bulk, and all of them degrade in a way that standard quality metrics do not detect: every individual item is fine, and the collection is redundant.

The natural formalization is a two-term objective. Let  $E$  be an embedding oracle and let  $X_n = \{x_1, \dots, x_n\}$  be the corpus so far. For a candidate  $c$ ,

$$J_\lambda(c \mid X_n) = \underbrace{\lambda \cdot \frac{1}{n} \sum_{i=1}^n \|E(c) - E(x_i)\|}_{\text{anchor: keep it typical}} - \underbrace{(1 - \lambda) \cdot \min_{i \leq n} \|E(c) - E(x_i)\|}_{\text{repulsion: keep it new}}$$

and we accept the candidate minimizing  $J_\lambda$ . The anchor term keeps generation on the manifold; the repulsion term pushes it away from what already exists. The interesting question is what happens to this as  $n \rightarrow \infty$ .

Our answer is that the objective is not where the difficulty lives. Both terms are cheap to maintain online and both behave predictably. The difficulty is that **the set of things the generator can propose is much smaller than it appears**, and the objective is powerless to enlarge it. Selection can only choose among what is offered. If the offering is confined to a thin slice, no selection rule — however clever, however well-tuned its  $\lambda$  — can produce diversity that the proposal distribution does not contain. Worse, as we show, aggressive selection under a fixed proposal distribution systematically degrades quality, because the only way to keep winning a novelty contest inside an exhausted region is to leave the manifold.

### I.1.1 The oracle asymmetry We assume throughout:

**Assumption (embedding oracle, no inverse).** We have access to  $E: \text{Text} \rightarrow \mathbb{R}^D$ , computable on demand. We have **no**  $E^{-1}: \mathbb{R}^D \rightarrow \text{Text}$ .

This is the structural fact that shapes the entire design, and it is worth dwelling on because it is easy to forget when writing down objectives. Given  $X_n$  we can compute, exactly and cheaply, the point in  $\mathbb{R}^D$  that would most improve any of our diversity measures — the direction of least occupied spectral energy, the centre of the largest empty ball, the point maximizing marginal coverage. Knowing that point is worth nothing on its own. There is no procedure that turns a target embedding back into a poem.

Consequently the system cannot *solve* for its next item. It can only:

1. **propose** — sample from  $p(\cdot \mid x)$  for some prompt  $x$  we are able to write;
2. **measure** — embed the proposals and score them;
3. **select** — keep one.

All of our leverage is therefore in step 1, in the choice of  $x$ , and that choice is expressible only in language. This is why the latent variables in our system are *language-valued* — named axes with named levels — rather than continuous codes: they are the only handles that reach the generator. It is also why the axis scoring of §4 exists. **The axis scoring is a surrogate for the missing inverse.** Since we cannot decode the direction we want to travel, we instead score the language-valued conditions we *can* write by how nearly their induced output distributions point that way.

### I.1.2 Contributions

- **A conditional-dimension theory of generative saturation** (§3). Five theorems relating the conditional support dimension  $m$  to the reachable manifold dimension  $d$ , each verified numerically. The practically important one is that saturation is governed by  $m$ , which is small, and not by  $d$ .
  - **An axis-scoring rule** (§4). A decision procedure for which latent variable to condition on next and which of its values to use, together with a principled termination/recursion signal. Includes an estimator bug we consider instructive enough to document (§4.3).
  - **A budget-matched evaluation protocol** (§5). Every method spends exactly 10,000 generator calls. We argue this is the only fair comparison for selection-based diversity methods, which buy diversity by discarding.
  - **Two domains with opposite diversity semantics** (§6): poetry (soft objective) and certification exam items (hard floor). The hard-floor domain exposes a failure — saturated banks silently fill with junk — that the soft domain hides.
  - **Live pilots** on `openai/gpt-5.6-luna` with locally-computed embeddings (§7).
- 

### I.2 Related work

**Diversity metrics.** The Vendi Score (Friedman & Dieng, 2022) is the exponential of the von Neumann entropy of a normalized similarity matrix, interpretable as an effective number of distinct items. We use it throughout, with a computational note: under a linear kernel on unit-normalized embeddings, the  $n \times n$  Gram matrix and the  $D \times D$  second-moment matrix share their nonzero spectrum, so the score costs  $O(nD^2 + D^3)$  rather than  $O(n^3)$ . We verify this to machine precision (max eigenvalue mismatch  $1.1 \times 10^{-16}$ ). This matters for our setting specifically: an infinite-horizon method needs an evaluation whose cost does not grow superlinearly in the thing it is evaluating.

**Selection and subset choice.** Determinantal point processes model repulsion via volume in feature space. Farthest-point /  $k$ -center greedy is the classical max-min packing heuristic and we use it as a baseline (`fps_post`) and, at the level of axis values, as a component (§4.4). SemDeDup-style near-duplicate removal at a fixed radius appears as `dedup_post`. All of these are *post hoc* selectors over a fixed pool, which is exactly their limitation in our setting: they cannot change what is in the pool.

**Quality-diversity.** MAP-Elites and its descendants maintain an archive indexed by hand-designed behavior descriptors, seeking an elite per cell. Our axis lattice is a close relative, with two differences that matter here: the descriptors are *elicited from the generator rather than designed by us*, and the lattice *refines itself* when a cell saturates, so the archive's resolution is not fixed in advance.

**Mode collapse in instruction-tuned models.** Repeated sampling from aligned models is well documented to produce low lexical and semantic entropy relative to base models. We take this as our empirical starting point and model it explicitly as skewed mixture weights over style modes.

**Novelty search.** Novelty search in evolutionary computation rewards behavioral distance from an archive. Our repulsion term is a direct analogue, and our §3.5 result is a formal statement of a hazard that community knows well in practice: unconstrained novelty pressure discovers degenerate behaviors, because degeneracy is genuinely novel.

**Coverage.** Part II treats the covering objective — maximizing the volume of a union of  $\epsilon$ -balls at a *finite* budget — which is submodular where max-min is not. §8 relates the two.

### I.3 Theory: conditional dimension governs saturation

**I.3.1 Setup** Let  $M \subseteq \mathbb{R}^D$  be the reachable semantic manifold,  $\dim M = d$ : the set of embeddings of texts the generator could produce under *some* prompt. For a fixed prompt  $x$ , let  $Sx \subseteq M$  be the support of  $E\#p(\cdot | x)$ , with  $\dim Sx = m$ .

The empirical claim behind everything below is that  $m \ll d$ . A prompt fixes topic, stance, form, register, and rhetorical strategy; what remains free is a low-dimensional residue — some lexical choice, some ordering, some surface variation. Our live pilot supports this directly: 60 poems generated from an identical naive prompt have a median nearest-neighbour cosine similarity of 0.911 and a Vendi Score of 2.55, meaning sixty samples behave like roughly two and a half distinct items.

Throughout,  $B(x, r)$  is the ball of radius  $r$  and we write  $gn$  for the min-gap of a fresh draw against a corpus of size  $n$ .

#### I.3.2 Fixed-prompt saturation

**Theorem 1.** Let  $Xn$  be  $n$  i.i.d. draws from a distribution with density bounded above and below on an  $m$ -dimensional manifold  $S$ . For a fresh draw  $c$  from the same distribution,  $\square [gn] = \Theta(n^{-1/m})$ .

*Proof sketch.*  $\square (gn > r) = \square (\text{no sample in } B(c, r)) = (1 - \mu(B(c, r)))^n$ , and  $\mu(B(c, r)) = \Theta(r^m)$  by density bounds on an  $m$ -manifold. So  $\square (gn > r) \approx \exp(-n\Theta(r^m))$ , which transitions at  $r = \Theta(n^{-1/m})$ ; integrating the tail gives the expectation.  $\square$

The content is the exponent. With  $d = 32$ , a naive reading predicts  $n^{-1/32}$  — essentially flat, novelty never exhausted. The true rate at  $m = 3$  is  $n^{-1/3}$ , which is roughly a  $10\times$  loss of headroom per  $1000\times$  of corpus. Measured slopes:  $-0.472$  ( $m = 2$ ),  $-0.309$  ( $m = 3$ ),  $-0.195$  ( $m = 5$ ), against  $-0.5$ ,  $-0.333$ ,  $-0.2$ . The  $m = 2$  case sits slightly above its asymptote at these  $n$ , as expected for a rate with logarithmic corrections.

**Corollary 1.1 (oversampling is not a strategy).** Drawing  $K$  candidates and keeping the most novel multiplies the expected gap by  $\Theta(K^{1/m})$  at best. To recover the headroom lost by a  $1000\times$  larger corpus,  $K$  must grow by 1000.

Verified at  $m = 3$ : best-of-8 yields a gain of 2.19 against  $K^{1/3} = 2.00$ . Selection is a constant-factor fix to an asymptotic problem.

#### I.3.3 The slice deficit

**Theorem 2.** If  $m < d$ , then  $\text{vold}(Sx) = 0$ , and the fraction of  $M$  within  $\varepsilon$  of  $Sx$  is  $\Theta(\varepsilon^{d-m})$ .

*Proof sketch.* The  $\varepsilon$ -neighbourhood of an  $m$ -dimensional set in  $d$  dimensions is a tube of volume  $\Theta(\varepsilon^{d-m}) \cdot \text{vol}_m(Sx)$  for  $\varepsilon$  below the reach of  $Sx$ . Dividing by  $\text{vold}(M)$  gives the claim.  $\square$

Measured at  $d = 8$ ,  $m = 2$ : fitted exponent 5.06 against a predicted 6, with covered fractions falling  $0.051 \square 0.0001$  as  $\varepsilon$  goes  $0.5 \square 0.12$ . The fitted exponent sits below prediction because at the larger  $\varepsilon$  the tube is no longer thin relative to the manifold, which flattens the low- $\varepsilon$  end of the fit; the qualitative claim — coverage collapsing polynomially in  $\varepsilon$  with a large exponent — is unambiguous.

The practical reading is severe. **A single prompt, sampled infinitely often, covers essentially none of what the model could write.** Not “less than we would like” — a fraction that goes to zero polynomially as the resolution of interest sharpens.

**I.3.4 Transversality: which prompt motion helps** Let a prompt schedule induce slices  $S1, \dots, SP$ . Write  $\text{Vocc}$  for the span of the corpus's dominant directions.

**Theorem 3.**  $\dim(\square j S_j) = \min(d, m + \text{rank}\{c_j - c_1\})$ , where  $c_j$  is the centre of  $S_j$ . In particular, displacements lying inside  $\text{span}(S_1)$  contribute nothing to the union's dimension.

*Proof sketch.* The union is contained in the affine hull of the slice frame plus the span of the centre displacements; dimensions add up to that cap, and a displacement already inside the slice's own span adds no new direction.  $\square$

Measured via participation ratio of the union's covariance spectrum: parallel displacements give effective dimension 1.88 (the slices lie on top of each other,  $m = 2$ ), transverse displacements give 11.18.




Figure 1. Measured saturation. (a) Expected min-gap of a fresh draw decays as  $n^{(-1/k)}$ , with fitted slopes matching theory for  $k = 3, 8, 32$ . (b) Best-of- $K$  oversampling buys only  $K^{1/k}$ .

This is the theorem that makes the axis scoring of §4 possible. It measures the value of conditioning on a new latent variable as **how much of its variation lies outside what the corpus already spans** — a computable quantity, in place of the intuition of how different it sounds.

**I.3.5 Novelty pressure leaves the manifold** Let the proposal distribution be a mixture  $(1 - p) \cdot \text{Pon} + p \cdot \text{Poff}$ , where  $\text{Poff}$  is a diffuse off-manifold component (incoherent text, degenerate output).

**Theorem 4.** Under best-of- $K$  max-min selection,  $\square$  (selected candidate is off-manifold)  $\square 1 - (1 - p)^K$ , i.e. the off-manifold candidate wins whenever one is present, for any corpus size. The selection rule is therefore inconsistent: it does not converge to sampling  $\text{Pon}$ .

*Proof sketch.*  $\text{Poff}$  has support far from the corpus, so its min-gap stochastically dominates that of any on-manifold candidate; the argmax picks it whenever it appears.  $\square$

Measured:  $\square$  (selected is off-manifold | one present) = 1.000 at both  $n = 50$  and  $n = 5000$ ; unconditional rate 0.0418 against a per-candidate rate of 0.005, an 8.4 $\times$  amplification at  $K = 8$ .

The consequence is not subtle in practice. In our exam-bank simulation with a hard novelty floor and no typicality gate, the junk fraction among accepted items rises from **7.9% in the first half of the bank to 48.2% in the second**. Once the legitimate item space is  $\delta$ -saturated, the only candidates still clearing the floor are the broken ones. **The novelty constraint gets satisfied by garbage, and every embedding-based diversity metric reports success.**

**Corollary 4.1.** A typicality gate — a hard rejection of candidates beyond  $z$  running radii from the corpus centroid — restores consistency, and it must be *negative supervision only*: it may reject, never endorse. Passing the gate is not evidence of quality.

**I.3.6 Breadth versus depth** Given budget  $B$ , split as  $P$  prompts  $\times n$  samples each, with per-prompt switching cost  $c$  (eliciting a spec, embedding it, occasionally refining), so  $B = P(c + n)$ .

**Theorem 5.** With  $c = 0$  the coverage-optimal depth is  $n^* = 1$ . For  $c > 0$  the optimum is interior and increases with  $c$ .

Measured: with free switching, depth 1 achieves coverage 0.995 while spending the entire budget on one prompt achieves 0.091 — a 10.9× difference. With  $c = 3$ ,  $n^* = 10$ ; with  $c = 30$ ,  $n^* = 30$ .




Figure 2. Breadth versus depth at fixed budget. With free prompt-switching the optimum is one sample per prompt; an interior optimum appears only once switching is priced.

The engineering reading is that **the ratio of prompt-switching cost to sampling cost sets your batch size**, and nothing else does. If constructing a new spec is as cheap as a generation, generate one item per spec.

**I.3.7 What this means for the objective** Returning to  $J\lambda$ : the anchor term is exactly computable from two running statistics. Writing  $\mu^*_{n^*}$  for the corpus centroid,

$$\frac{1}{n} \sum_i \|c - x_i\|^2 = \|c - \mu_n\|^2 + \frac{1}{n} \sum_i \|x_i - \mu_n\|^2$$

verified to zero relative error. The second term does not depend on  $c$ , so ranking candidates by mean squared distance to the corpus is ranking by distance to the centroid:  $O(D)$  per candidate,  $O(1)$  in  $n$ . The repulsion term needs a nearest-neighbour query, which we bound by subsampling. Neither term is the bottleneck. **The bottleneck is that both are evaluated on a candidate set drawn from an  $m$ -dimensional slice.**

## I.4 Scoring candidate axes

**I.4.1 The question the axis scoring answers** Theorem 3 says progress requires moving the prompt transversely to what the corpus already spans. Theorem 5 says move often. Together they pose a concrete question at every step: **among the latent variables we could condition on next, which one moves the slice most transversely, per unit of budget?** Since we have no inverse oracle, this is the closest available substitute for computing the ideal next item.

Each candidate axis  $a$  has language-valued levels  $L(a)$ . We embed the level descriptions and read off four quantities.

**I.4.2 The four factors** **Spread** — mean pairwise distance between  $a$ 's level embeddings. Are this axis's values actually different *from each other*? An axis whose levels are near-synonyms cannot separate outputs however it is sampled. This catches the most common failure of LLM-proposed axes: plausible-sounding dimensions whose levels all mean the same thing.

**Transversality** — the fraction of  $a$ 's level-to-level variation lying outside the corpus's occupied eigenspace (top eigenvectors of the second-moment matrix carrying 80% of spectral energy). This is Theorem 3 made into a number. An axis whose levels differ only along directions the corpus already spans adds density, not reach.

**Independence** —  $1 - \max$  principal-angle cosine between  $a$ 's variation subspace and those of the axes **already in use**. An axis that restates an active axis is redundant however transverse it looks alone.

**Headroom** — a normalized entropy deficit of how unevenly the corpus has sampled  $a$ 's levels, plus the fraction of levels never used. This is the only time-varying factor; it is what makes the ranking change as generation proceeds rather than being a fixed property of the axis set. An axis can be excellent and already spent.

The **promise** of an axis is the product:

$$\text{promise}(a) = \text{spread}(a) \cdot \text{transversality}(a) \cdot \text{independence}(a) \cdot \text{headroom}(a)$$

Multiplicative, not additive, and deliberately so: an axis needs all four, and any one at zero should zero the score rather than be averaged away by the others. A high-spread axis with zero transversality is a distinction without a difference; a perfect axis with zero headroom is already spent.

**I.4.3 An instructive bug** Our first implementation scored independence against *all other candidate axes*. On a ground-truth test with four synthetic axes — one redundant-in-span, one with near-synonym levels, one genuinely new, and an exact duplicate of the genuinely-new one — the two excellent duplicates each drove the other's independence to zero, and **both scored below the useless near-synonym axis** (promise  $5 \times 10^{-5}$  versus 0.020).

The fix is conceptual rather than numerical. Independence is a property of a candidate *relative to what is already in use*, never relative to its rivals. Redundancy *among candidates* is a selection problem, handled by greedy `select_axis_set`, which re-scores after each pick so a duplicate is penalized at the point where the redundancy actually comes into existence. After the fix the ranking is correct (genuinely-new 1.314, its duplicate 1.313, near-synonym 0.111, redundant-in-span 0.003) and greedy selection of two axes never takes both an axis and its duplicate.

We report this because the same error is easy to make in any diversity-scoring system, and because it fails in a direction that looks reasonable: it *punishes* exactly the axes worth keeping, in proportion to how well they were proposed.

**I.4.4 Choosing values, and recursing** Once an axis is chosen we do not sample its levels uniformly. We take the **max-min subset** of its level embeddings by greedy farthest-point search, seeded deterministically at the level furthest from the level centroid — the packing problem again, one level down.

When the best available promise falls below a floor, the axis scoring emits decision: "refine" instead of "condition". This is the **exhaustion signal**, and it is the honest one: it says the current lattice has nothing transverse left to offer, and no amount of further sampling will change that. The system then asks the generator to split the exhausted level into finer sub-levels (§5.2), producing a child axis scored by this same axis scoring. On the ground-truth test with only in-span axes available, the axis scoring correctly returns refine.

Cost is  $O(|A| \cdot D^2)$  per decision and does not grow with  $n$ .

---

## I.5 Method

**I.5.1 The stack** Per accepted item, a bounded number of model calls regardless of corpus size:

1. **Spec choice.** Sample a pool of specs from the current axis lattice; embed their descriptions; keep the one with the largest residual outside the corpus's occupied eigenspace. Orthogonalization at the level of *conditions*, not outputs.
2. **Generation.**  $K$  parallel completions conditioned on the spec, with the spec's levels stated as **contracts** — behaviors the text must exhibit, not suggestions — plus an explicit avoid-list from the ledger.
3. **Judging.** One separate call scores craft (or, for exam items, validity) and checks each required behavior. Generation never grades itself.
4. **Selection.** Utility =  $0.4 \cdot \text{quality} + 0.35 \cdot \text{orthogonality} + 0.25 \cdot \text{capped gap}$ , behind a typicality gate. The gap term is **capped** at 2 running scales — this is what prevents the Theorem 4 pathology, since an uncapped gap reward is unbounded and off-manifold candidates always win it.
5. **Mining.** Every 25 accepts, show the model 50 sampled items and ask what they have in common. Append the answer to a JSONL ledger; a bounded slice becomes the next prompts' avoid-list.
6. **Refinement.** On a saturation signal, ask the generator to split the dominant saturated cell.

Every state the policy reads is bounded: running centroid and second-moment, EMA scales, a fixed ledger slice, a recent window plus a bounded random subsample of older items. The marginal cost of item 10,000 equals that of item 100.

**I.5.2 Eliciting and refining latent behaviors** We ask the generator for axes that are near-orthogonal, structural rather than topical, and whose levels are all usable — explicitly steering away from subject matter, length, and rhyme toward craft decisions. On the live run this produced axes including *temporal stance*, *relationship to its own claim*, *syntactic weather*, and *the register the poem refuses*.

Refinement shows the model the saturated cell and the mined attractors and asks it either to split a level into finer sub-levels or to mint a new axis meaningful only inside that cell. Refinement axes are **conditional** — they apply only when their parent level was chosen — so the lattice is a tree, and refining a saturated region does not inflate cost elsewhere.

**I.5.3 Why "orthogonalize", not "randomize"** Forcing randomness — raising temperature, injecting random seed words — buys variance in the surface while leaving the mode structure intact, and degrades quality monotonically because temperature cannot distinguish *surprising* from *wrong*. Our budget-matched results bear this out: `high_temp` achieves the highest raw Vendi Score of any method (28.12) at 4.6× the junk rate of naive sampling, and no quality gain whatsoever (0.671 versus 0.669).

Forcing approximate orthogonality asks a different question — which direction is this corpus not yet spending energy on — which has a computable answer, improves rather than degrades quality when paired with a quality term, and remains meaningful as  $n$  grows. "Be random" gets no harder to satisfy and no more useful.

---

## I.6 Experiments

All simulations are seeded and reproducible. Every mathematical claim is checked numerically; a failure there invalidates the corresponding claim here.

**I.6.1 The conditional world: the central result** We simulate a world where a spec genuinely selects a slice: ambient  $\mathbb{R}^{64}$ , reachable manifold  $d = 24$ , conditional dimension  $m = 3$ , additive composition over 5 axes  $\times$  4 levels, rare off-manifold junk. Budget 10,000 generations,  $K = 8$ .

| policy                           | Vendi        | quality | reach dim | lattice | inability    |
|----------------------------------|--------------|---------|-----------|---------|--------------|
| fixed prompt                     | <b>1.06</b>  | 0.668   | 13        | 1.0e3   | -0.104       |
| random spec                      | 6.70         | 0.724   | 13        | 1.0e3   | 0.867        |
| axis scoring spec                | 7.93         | 0.716   | 13        | 1.0e3   | 0.817        |
| random spec + refinement         | 7.33         | 0.727   | <b>24</b> | 3.3e4   | 0.681        |
| axis scoring + refinement        | 10.60        | 0.740   | <b>24</b> | 2.9e4   | <b>0.598</b> |
| axis scoring + guided refinement | <b>11.23</b> | 0.722   | <b>24</b> | 2.9e4   | 0.696        |

A fixed prompt yields a Vendi Score of **1.06** from 1,250 accepted items: the corpus is effectively one item. The full stack reaches **11.23**, a 10.6× gain at identical budget. Refinement lifts reachable dimension from 13 to 24 — Theorem 3 in action — while the lattice grows 1e3  $\square$  3e4.

Conditional world: a prompt selects an  $m$ -dimensional slice ( $d = 24, m = 3$ )




Figure 3. The conditional world ( $d = 24, m = 3$ ). (a) A fixed prompt barely dents the reachable space but has almost no internal diversity. (b) Only recursive refinement raises the reachable-dimension ceiling. (c) Achieved diversity: 1.06 for a fixed prompt against 11.23 for the full stack.

Refinement direction matters: children displaced isotropically from saturated modes reduce self-headroom (1.20 versus 1.38 unrefined), since tighter spread near the parent concentrates mass where the corpus already sits. Refinement pays only along directions transverse to the occupied span, which is what the transversality term selects (Vendi 11.23 versus 10.60 for unguided refinement).

Note that lattice size and reachable dimension **come apart**. Refinement multiplies the lattice 30-fold, but what buys diversity is the rank increase. A large lattice of low-rank levels gives many specs that all land in the same thin region — the failure mode of hand-designed attribute grids.

**A trap worth naming.** The fixed prompt has the *best* inability score (−0.104, essentially zero) while having the worst diversity. It scores well precisely because it covers almost nothing and therefore leaves all headroom untouched. Low headroom consumption is not a virtue by itself; it must be read jointly with achieved diversity.

**I.6.2 Budget-matched comparison and ablations** Every method spends exactly 10,000 generations on one world. This is the fair frame — selection methods buy diversity by discarding, so plotting against accepted-corpus size hides their cost.

| method                  | $n$ kept | Vendi | quality | junk  | median gap |
|-------------------------|----------|-------|---------|-------|------------|
| naive                   | 10000    | 21.29 | 0.669   | 0.004 | 0.966      |
| high temperature        | 10000    | 28.12 | 0.671   | 0.020 | 1.542      |
| random latent / persona | 10000    | 26.76 | 0.700   | 0.005 | 0.936      |
| dedup (SemDeDup-style)  | 9517     | 21.56 | 0.671   | 0.005 | 0.978      |
| farthest-point post hoc | 1250     | 22.63 | 0.622   | 0.041 | 1.445      |

| method                 | $n$ kept | Vendi        | quality      | junk         | median gap |
|------------------------|----------|--------------|--------------|--------------|------------|
| online max-min         | 1250     | 22.48        | 0.622        | 0.044        | 1.359      |
| <b>IHD (no ledger)</b> | 1250     | <b>23.80</b> | <b>0.750</b> | <b>0.001</b> | 1.170      |
| IHD full (with ledger) | 312      | 22.60        | 0.763        | 0.000        | 1.268      |
| — no gate              | 312      | 23.00        | 0.766        | 0.010        | 1.254      |
| — no orthogonality     | 312      | 21.65        | 0.759        | 0.000        | 1.227      |
| — no quality term      | 312      | 24.35        | 0.664        | 0.013        | 1.389      |
| — no repulsion         | 312      | 21.44        | 0.783        | 0.000        | 1.101      |

Reading the ablations: removing **orthogonality** costs 2.15 Vendi ( $23.80 - 21.65$ ), the largest diversity loss, confirming it as the load-bearing diversity term. Removing **quality** raises Vendi to 24.35 while dropping quality  $0.750 - 0.664$  and raising junk  $13\times$ , which is the expected trade and the reason the term is there. Removing the **gate** raises junk from 0.001 to 0.010. Removing **repulsion** gives the highest quality (0.783) and the lowest Vendi — it collapses toward the safe centre of the distribution.




Figure 4. Budget-matched comparison and ablations at  $B = 10,000$  generations. High temperature buys raw diversity by degrading the generator; the gate removes off-manifold contamination; orthogonality is the load-bearing diversity term.

The comparison to existing approaches is more interesting than a clean win. `high_temp` attains the highest Vendi of any method. We do not regard this as it beating our method, and the reason is visible in the other columns: it achieves it with  $20\times$  the junk rate and no quality gain. **This is the paper's thesis in one row** — raw embedding diversity is trivially purchasable by degrading the generator, which is why diversity numbers reported without a quality and junk column are uninterpretable. Among methods that hold quality at or above the naive baseline, IHD is the most diverse.

**An honest cost finding.** IHD-with-ledger keeps only 312 items from the same budget, because our simulated ledger is implemented as  $4\times$  rejection oversampling, and it buys no Vendi over the ledger-free variant. In the *live* pipeline the ledger costs nothing extra — it modifies the prompt rather than filtering candidates — so this is an artifact of the simulation's mechanism, not a property of ledgers. We report it because the simulated version is what we measured.

**I.6.3 Poetry world at  $n = 10,000$**  Six acquisition policies (figure 1). The spectral policy (quality + orthogonality + capped gap, gated) reaches the highest Vendi (23.28) with **zero** junk in the last 2,000 items and the highest quality (0.753), while pure max-min holds a much larger min-gap (2.26 versus  $\sim 1.15$ ) at 3.9% junk and materially lower quality (0.626). Ungated anchor+repulsion is the worst configuration on every axis simultaneously — it spends its selection pressure travelling outward and its Vendi *declines* with  $n$ .

**I.6.4 Exam items: a hard floor and finite packing** Exam items invert the semantics. Two operational items closer than  $\delta$  are “enemy items” — seeing one answers the other — which is a test-security failure at *any* bank size. The floor is a constraint, never a term in a weighted sum. Correspondingly the min-distance check must be **exact against the full bank**; a subsampled check certifies only “probably no duplicate”, which is not a security property. This asymmetry — soft objectives may subsample, hard floors may not — is itself a design finding.

Item templates expose few manipulable slots, so each mode is a low-intrinsic-dimension disk and the  $\delta$ -packing number is finite and small. Our bound  $N_{\max} \leq M(2\text{Reff}/\delta + 1)^k$  gives 5,543 for our parameters; the largest gated bank reaches 1,464 real items, respecting it.

The important result is what happens without a gate. Ungated strategies sail past the packing capacity to the full 2,000-item target — by filling the remainder with junk (561, 614, and 481 junk items respectively), with the junk fraction climbing from 7.9% to 48.2% across the bank. **The hard novelty constraint was being satisfied by broken items.** Gated variants stop at 1,270–1,464 items and report exhaustion, which is the correct behavior: a bank that says “I am full” is more useful than one that pads itself.

Cost per accepted item rises from  $\sim 14$  generations early to  $\sim 40$  at saturation, and the gated runs pay 135–207 generations per item overall.

**I.6.5 The coverage horizon,  $5 \rightarrow 10,000$**  We probe, at log-spaced checkpoints, the best min-gap a *fresh coherent* candidate can achieve against the corpus so far — the novelty headroom remaining — and normalizes it into an *inability to be novel*,  $1 - h(g)/h(5)$ . Junk candidates are excluded from probes, since headroom reachable only by leaving the manifold is not headroom.

Every fixed-generator method converges toward exhaustion: naive 53.3%, random-latent 49.0%, online max-min 55.7%, IHD selection 48.1%. High temperature appears best (40.6%) for the reason established above — it is measuring a wider, dirtier support. **No selection policy escapes the horizon.** Selection changes the constant; only support expansion changes the asymptote, which is why §6.1’s refinement result is the one that matters.

Coverage horizon: 5  $\rightarrow$  10,000 generations




Figure 5. The coverage horizon closing from 5 to 10,000 generations. Every fixed-generator policy converges toward exhaustion.

## I.7 The real experiment: 43,000 generations, three domains, seven methods

Everything to this point is either a theorem or a simulation. This section is the measurement. Generator: openai/gpt-5.6-luna via OpenRouter. Text embeddings: nomic-embed-text (768-d) served locally by Ol-

lama, so embedding is free and the loop is never rate-limited by its own measuring instrument. Images: gpt-image-1-mini. Image embeddings: CLIP ViT-B/32. Audio: Lyria 3 Pro instrumentals. Audio embeddings: CLAP and MERT. Every corpus is append-only JSONL with embeddings checkpointed alongside, resumable after a kill.

**I.7.1 Domains DALL-E instructions for post-modern artworks.** Diversity *is* the product: a clustered instruction set renders to a clustered image set. Chosen also because it lets us measure the same corpus in three spaces — words, text-embedding, and the rendered picture.

**Psychometric test questions.** Multiple-choice items assessing general reasoning. Two items with the same construct and the same surface are redundant at best and, in an operational bank, a validity problem.

**Instrumental-music prompts.** The longest chain in the paper: latent axes  $\square$  a text prompt  $\square$  a two-minute instrumental  $\square$  an audio embedding. We render with Lyria 3 Pro (lyria-3-pro-preview), which returns roughly two minutes of audio per call in about 20 seconds, and — usefully — also emits its own section plan ([A0]) ([B1]) ([C2]) ([D3]), a model-reported description of the form it chose that gives us a discrete structural signal without any waveform analysis.

**A prompt-length decision that is really a methodological one.** Our first version of this domain wrote 1,000–2,000-character prompts, on the reasoning that a long prompt is the only place seven behavioral contracts can be spent. That reasoning is wrong, and wrong in a way this paper should be the first to catch. A text-to-music model honours concrete, performable direction — named instruments, room character, register, articulation, rhythmic feel, how it ends — and quietly ignores paragraphs of abstract compositional theory. Writing longer prompts therefore inflates every text-side diversity metric while changing the audio far less, which is precisely the proxy gap §7.6 measures. Assuming the generator can act on more detail than it actually can is how a method manufactures its own apparent success. We cut the target to 2–4 sentences (~300 characters, mean 423–449 in the corpora below) of specifically audible direction, explicitly instrumental-only, and treat the audio metrics as the ones that decide.

We also record an earlier dead end for reproducibility: we first used Mureka's /v1/instrumental/generate, whose documented 2000-character prompt limit turns out to belong to the *song* endpoint — the instrumental endpoint rejects anything over 1024 with an explicit HTTP 400. We found that against the live API rather than the docs. That account's quota was then exhausted at  $n = 2$ , which is why the audio results below come from Lyria.

**I.7.2 Methods compared** Five of these are real published approaches, implemented faithfully rather than as strawmen, each getting the same generator, the same embedder, and the same budget accounting as ours.

| arm           | what it is                                                                                                         |
|---------------|--------------------------------------------------------------------------------------------------------------------|
| naive         | plain repeated prompting — the honest floor                                                                        |
| high_temp     | temperature 1.6 — the trivial diversity lever                                                                      |
| self_instruct | <b>Self-Instruct / Alpaca:</b> few-shot exemplars sampled from the pool, plus ROUGE-L $\leq 0.7$ rejection against |
| evol_instruct | <b>Evol-Instruct (WizardLM):</b> sample an existing item, apply a random evolution operator (deepen, concret       |
| persona       | <b>Persona-Hub / AttrPrompt:</b> a flat catalogue of personas $\times$ attributes, sampled uniformly               |
| ihd           | ours: recursively elicited axes, spec-level orthogonalization, capped repulsion behind a typicality gate, ap       |
| vision        | ours, plus steering on the <i>rendered image</i> (§7.6)                                                            |

persona is the load-bearing comparison. It has language-valued latent conditioning — the same basic idea as ours — but no orthogonality selection, no ledger, and no recursive refinement. The gap between persona and ihd isolates exactly what those three components buy.

One faithfulness caveat we chose deliberately: Self-Instruct's ROUGE filter compares a candidate against the entire pool, which is  $O(n)$  longest-common-subsequence computations per candidate and comes to dominate the loop at  $n$  in the thousands. We compare against a bounded random sample of 120 pool members.

This makes our reimplementation *weaker* than the original at large  $n$ , and that is itself the point: the original's redundancy check does not have an infinite horizon, because its cost grows linearly in the corpus it is protecting.

### I.7.3 Mode collapse is not a metaphor Two 10,000-item corpora, one call each, no selection:

| domain              | $n$    | unique texts | exact-duplicate rate | most-repeated item |
|---------------------|--------|--------------|----------------------|--------------------|
| DALL·E instructions | 10,000 | 10,000       | 0.000                | 1×                 |
| psychometric items  | 10,000 | ~1,656       | <b>0.723</b>         | <b>1,637×</b>      |

In the psychometric domain, 72.3% of a ten-thousand-item corpus is exact duplicate text, and one single question — “What number comes next in the sequence: 2, 6, 12, 20, 30, ?” — accounts for 1,637 of them. No embedding, no threshold, and no interpretation is required to see this; it is byte-identical repetition, and it is what a strong instruction-tuned model does when asked the same reasonable question ten thousand times.

Two things make this more than a curiosity. First, **temperature barely helps**: at  $T = 1.6$  the duplicate rate is still 0.616. Second, **conditioning nearly eliminates it**: persona conditioning drops it to 0.002. That single contrast is the paper's thesis reduced to two numbers — randomness does not buy diversity, conditioning does.

It also shows the failure is domain-shaped and invisible from one vantage point. The identical pipeline, prompt style, and model produce zero duplicates on DALL·E instructions. A practitioner who validated their pipeline on the first domain and deployed it on the second would ship a bank that is three-quarters one question.




Figure 6. Real gpt-5.6-luna corpora,  $n = 5$  to 10,000. Literal measures (top row) and latent measures (bottom row) do not agree about what is happening.

### I.7.4 Literal and latent diversity move in opposite directions Measuring the DALL·E naive corpus as it grows from $n = 5$ to $n = 10,000$ , in both spaces:

| $n$   | distinct-2 | 4-gram self-repetition | $n$ -gram Vendi | centered embedding Vendi |
|-------|------------|------------------------|-----------------|--------------------------|
| 5     | 0.870      | 0.031                  | 4.9             | 3.79                     |
| 40    | 0.492      | 0.211                  | 30.2            | 24.13                    |
| 300   | 0.252      | 0.384                  | 142.5           | 55.70                    |
| 1,000 | 0.145      | 0.520                  | 246.6           | 63.05                    |
| 1,750 | 0.112      | 0.577                  | 283.3           | 65.30                    |

Literal diversity collapses monotonically — by  $n = 1,750$  more than half of each new instruction's 4-grams have already appeared — while latent diversity *rises* and then saturates. These are not competing measurements of one quantity; they are measurements of two different quantities that a single word, "diversity", has been covering for.

**A methodological correction we owe the reader.** The uncentered embedding Vendi on this corpus reads  $1.63 \square 2.33$ , which would suggest ten thousand instructions behave like two distinct items. That number is mostly an artifact of the kernel. Same-domain embeddings sit in a narrow cone — mean pairwise cosine similarity is 0.883 here and 0.788 on the psychometric corpus — so the Gram spectrum is dominated by the shared mean direction and the score compresses toward 1. After removing the mean direction the same corpus reads  $3.79 \square 65.30$ , which is the honest curve and the one in the table. We report both, and we suggest any embedding-based diversity result publish the corpus's pairwise-similarity distribution alongside it, because the number is meaningless without it. Part II measured a mean pairwise cosine of 0.444 on its own pool with the same embedder — the cone is corpus-dependent, not a fixed property of the embedder, which is exactly why it has to be reported rather than assumed.




Figure 7. The decoupling, normalised to each series' value at  $n = 5$ : literal diversity falls while latent diversity rises.

**I.7.5 Competitive comparison at matched  $n$**  All seven arms, real corpora, DALL-E domain, every arm evaluated on the same number of accepted items:

| arm                  | distinct-2 $\square$ | self-repetition $\square$ | $n$ -gram Vendi $\square$ | centered Vendi $\square$ | median NN distance $\square$ |
|----------------------|----------------------|---------------------------|---------------------------|--------------------------|------------------------------|
| naive                | 0.334                | 0.309                     | 80.6                      | 43.99                    | 0.046                        |
| high temperature     | 0.334                | 0.301                     | 82.7                      | 44.48                    | 0.052                        |
| Evol-Instruct        | 0.413                | 0.327                     | 81.9                      | 43.26                    | <b>0.025</b>                 |
| Self-Instruct        | 0.558                | 0.137                     | 108.8                     | 53.05                    | 0.115                        |
| persona conditioning | 0.619                | 0.121                     | 101.0                     | 49.45                    | 0.085                        |
| <b>IHD (ours)</b>    | 0.703                | 0.056                     | 122.7                     | 65.47                    | 0.179                        |

| arm                 | distinct-2 $\square$ | self-repetition $\square$ | $n$ -gram Vendi $\square$ | centered Vendi $\square$ | median NN distance $\square$ |
|---------------------|----------------------|---------------------------|---------------------------|--------------------------|------------------------------|
| IHD + vision (ours) | 0.720                | 0.038                     | 124.9                     | 73.25                    | 0.187                        |

Ours wins every column, and adding vision steering improves on text-only IHD by a further 12% in centered Vendi. Three of the baselines fail in ways worth naming.

**High temperature buys nothing.** distinct-2 of 0.3342 against naive's 0.3341. Turning up the temperature moved the third decimal place. In the simulated world (§6.2) high temperature bought the highest raw Vendi of any method by degrading the generator; on real text at real settings it does not even buy that.

**Evol-Instruct is worse than naive at the thing diversity methods exist for.** Its median nearest-neighbour distance is 0.025, against naive's 0.046 — it produces items *closer* to the existing corpus than independent sampling does. The mechanism is working as designed: mutating an existing item produces a neighbour of that item. Evol-Instruct raises *distinct-2* (0.413 vs 0.334) because the mutations reword things, so a purely literal evaluation would score it as an improvement. Measuring both levels is what catches it.

**Self-Instruct's guard is literal, so it defends the literal level only.** It achieves the second-best self-repetition (0.137) — its ROUGE filter is doing real work — while its centered Vendi (53.05) trails ours by 20 points. A ROUGE-L threshold cannot see semantic redundancy, and semantic redundancy is what remains once the lexical kind is filtered.




Figure 8. Competitive comparison on real corpora at matched  $n$ , both domains, across four metrics.

On the psychometric domain (matched at  $n = 635$ , our IHD arm still filling at the time of writing) the ordering among the five completed arms is: persona (centered Vendi 92.4, dedup'd 92.6) > Evol-Instruct (61.1) > Self-Instruct (54.2)  $\square$  high temperature (7.0)  $\approx$  naive (7.1). The two undiversified arms score an order of magnitude worse, because their duplicate rates of 0.62–0.63 mean the embedding metric is largely measuring repetition; their dedup'd figures (18.3 and 16.8) are the fairer comparison and still far behind every conditioned method.

**I.7.6 The proxy problem: text diversity is nearly blind to image diversity** We rendered the first 199 DALL-E instructions from the naive corpus to actual images and embedded them with CLIP. The result is the most consequential measurement in this section:

Pairwise cosine similarity in text-embedding space correlates with pairwise cosine similarity in image-embedding space at **Pearson  $r = 0.155$** .

Knowing that two instructions are semantically far apart tells you very nearly nothing about whether the two pictures look different. Every text-side diversity method in the table above — ours included — is optimizing a proxy that explains roughly 2% of the variance in the thing the user actually receives.

The qualitative version is more damning than the correlation. Six independently generated instructions, from a corpus with a 0.000 exact-duplicate rate and healthy lexical diversity, render to six near-interchangeable pictures: the same magenta-and-cyan palette, a classical marble bust, neon signage, halftone collage, a receding grid. A vision judge shown samples of the set rates its distinctness 6–7 out of 10 and names the attractors precisely — *"muted beige, cream, brown, ochre and black foundations accented by saturated cyan/teal, turquoise, pink"*, *"appropriation of canonical or religious imagery, especially Mona Lisa-like female portraits"*, *"frontal, museum-like presentation with centered, symmetrical compositions"*.

This is a second mode collapse, downstream of ours, contributed by the image model and by the fact that much of what varies in the text ("post-modern", "appropriated source") lands in the same visual place.

Text vs vision diversity on 199 rendered post-modern artworks (gpt-image-1-mini + CLIP)




Figure 9. Text versus vision diversity on 199 rendered artworks. (c) Pairwise similarities in the two spaces correlate at only  $r = 0.155$ .

First 48 rendered instructions (naive prompting, gpt-image-1-mini)




Figure 10. The first 48 rendered instructions. The corpus has a 0.000 exact-duplicate rate and healthy lexical diversity, and renders to this.

**The vision-steered arm** closes the loop where the product actually lives. It renders a bounded sample of accepted instructions, embeds them with CLIP, and feeds two things back into the text-side loop: a least-squares map from the crowded *image* directions into instruction-embedding space, so the text-side orthogonality term can push away from visual redundancy it cannot itself perceive; and mined *visual* attractors from the vision judge, appended to the same append-only ledger as the textual ones and repelled against in subsequent prompts. It is the paper's mechanism applied one level down: the ledger already repels against what the model keeps saying, and now also against what it keeps showing. It is the best arm in the table.

**I.7.7 Attractor mining works, and is legible** The mining step produces findings specific enough to act on. From the poetry pilot, round 1 ( $n = 25$ ): *"self-correction and immediate retraction: speakers repeatedly interrupt their own claims"; "ceremonial or institutional address: the voice of a bell-ringer, town crier, registrar, witness"; "threshold imagery and delayed passage: doors, gates, windows, bridges, shores"*. By round 2 ( $n = 50$ ) it tracks the corpus's *drift* rather than restating round 1, noting that technical vocabulary is now being placed inside mythic frames and that recursive epistemic backtracking has become structural.

Two observations. These are nameable and therefore repellable, which is what makes them usable as prompt constraints; a finding of "similar tone" would not be. And several of them are artifacts of *our own axis elicitation* — asking for craft-level axes like "relationship to its own claim" reliably produces self-correcting speakers. **The system's own conditioning becomes the next attractor.** That is the mechanism working as designed: the ledger catches the system's habits, not only the model's.

**I.7.8 Cost, and what did not finish** The text corpora comprise 43,171 real generations for roughly \$8.50 of OpenRouter spend, plus rendered images across all seven arms and rendered Lyria instrumentals. Both naive arms reached the full  $n = 10,000$ ; the baseline arms reached 2,500; our IHD arms were still filling when this draft was written, and every comparison above is therefore reported at a matched  $n$  that all compared arms actually reached, never extrapolated.

**7.9a The third domain: instrumental-music prompts** The music domain replicates the result on a third, structurally different artifact. The generator proposed seven compositional axes — *formal trajectory, inter-layer rhythmic relationship, harmonic motion, timbral centre of gravity, density contour, pulse relationship, opening and closing frame* — which are craft decisions rather than genre labels, and each of which can be expressed as something audible rather than as theory.

The prompt-level table below is the earlier, long-prompt version of the corpus and is reported for the text-side comparison only; the audio measurements in §7.9b use the shorter, performable prompts described in §7.1.

| arm               | $n$ | distinct-2 $\square$ | self-repetition $\square$ | $n$ -gram Vendi $\square$ | centered Vendi $\square$ | median NN dist $\square$ |
|-------------------|-----|----------------------|---------------------------|---------------------------|--------------------------|--------------------------|
| naive             | 100 | 0.348                | 0.291                     | 56.0                      | 22.99                    | 0.027                    |
| <b>IHD (ours)</b> | 100 | <b>0.545</b>         | <b>0.140</b>              | <b>75.3</b>               | <b>43.03</b>             | <b>0.079</b>             |

A 1.87 $\times$  gain in centered Vendi, self-repetition halved, and roughly three times the room between nearest neighbours — at matched  $n$ , matched prompt length, and the same generator. Neither arm produces exact duplicates, so this is a purely semantic effect. Three domains, three replications, and in each the mechanism that moves the number is conditioning rather than sampling temperature.

The audio domain is the least complete. The pipeline runs end-to-end — prompts generated, instrumentals rendered, CLAP and MERT embeddings computed with a mean/spread/temporal-difference aggregate over windows so that two tracks with identical average timbre but different *form* do not collide — but at a scale that supports description, not inference. We report it as an existence proof and a set of measurement machinery, not as a result.

**I.7.9 Head-to-head against a human-written exam bank** The comparisons above are against our own reimplementations, which is the right experiment for isolating mechanisms and the wrong one for the question *is this actually good?* — a reimplementation can be weak in ways that flatter us. So we also compare against an item bank that humans wrote: **MMLU**, ~14,000 multiple-choice items drawn from real practice exams and textbooks across 57 subjects, by many authors, with editorial review, over years. All arms are sampled uniformly at random and evaluated at matched  $n = 1,000$  with identical metric code.

| source                      | exact-dup $\square$ | distinct-2 $\square$ | self-repetition $\square$ | $n$ -gram Vendi $\square$ | centered Vendi $\square$ | med         |
|-----------------------------|---------------------|----------------------|---------------------------|---------------------------|--------------------------|-------------|
| <b>MMLU (human-written)</b> | 0.0040              | <b>0.5888</b>        | 0.0585                    | 495.2                     | <b>197.50</b>            | <b>0.28</b> |
| naive gpt-5.6-luna          | 0.6560              | 0.1196               | 0.8383                    | 12.1                      | 6.52                     | 0.00        |
| high temperature            | 0.6840              | 0.1161               | 0.8567                    | 10.2                      | 6.66                     | 0.00        |
| Self-Instruct               | 0.0280              | 0.1990               | 0.4797                    | 277.6                     | 57.22                    | 0.03        |
| Evol-Instruct               | 0.0060              | 0.2642               | 0.3855                    | 381.9                     | 89.23                    | 0.07        |
| persona conditioning        | 0.0040              | 0.3887               | 0.1711                    | 543.3                     | 99.71                    | 0.13        |
| <b>IHD (ours)</b>           | <b>0.0000</b>       | 0.4377               | <b>0.0503</b>             | <b>622.2</b>              | 135.58                   | 0.22        |

Two results, and we want to be precise about which is which.

**Against every synthetic method, we win on every column.** Centered Vendi 135.6 against the best baseline's 99.7 — a 36% margin over persona conditioning, which is the closest published relative of our approach and differs from it exactly by the three components we add. Self-repetition 0.050 against persona's 0.171. The

undiversified arms are not close: naive and high-temperature prompting score  $\sim 6.5$  centered Vendi, because two-thirds of their corpora are byte-identical duplicates, and their median nearest-neighbour distance is exactly 0.0000 — the median item has a perfect twin.

**Against the human bank, we split.** We *beat* MMLU on exact duplication (0.0000 vs 0.0040 — the human bank has some), on 4-gram self-repetition (0.050 vs 0.059), and on  $n$ -gram Vendi (622 vs 495): our items reuse less language than human-written ones do. We *lose* on the semantic measures — centered Vendi 135.6 against 197.5, and median nearest-neighbour distance 0.224 against 0.284. We reach roughly 69% of a human exam bank's effective semantic diversity.

That gap is the honest measure of what is left. MMLU's spread comes from 57 genuinely different subjects; ours comes from an axis lattice a single model proposed in one call, refined a handful of times. The lexical result says our surface variety already exceeds human-authored items; the semantic result says our *conceptual* variety does not, and that closing the remaining 31% is a question about how much genuinely different subject matter the generator can be induced to reach — which is precisely the reachable-dimension question of §3, not a tuning problem.

**I.7.10 A circularity in our own favour, and the metrics that are free of it** Our selection rule maximizes a weighted sum of embedding-space orthogonality and embedding-space min-gap. We then report embedding-space diversity metrics. Those two facts are not independent: the centered Vendi Score is a monotone function of how flat the embedding Gram spectrum is, which is close to exactly what the orthogonality term climbs, and the median nearest-neighbour distance *is* the min-gap term. To that extent, our wins on `embed_vendi_centered` and `median_nn_cos_dist` are partly tautological — we optimized them, and the baselines did not.

We flag this rather than let it pass, because Part II found a sharper version of the same error in its own benchmark (a selector scored against the very reference set it had optimized against) and had to revise its headline number downward after fixing it.

The defence is that our method never observes the *literal* metrics at all. It reads embeddings; it has no access to token counts,  $n$ -gram overlap, or string identity. So exact-duplicate rate, distinct-2, 4-gram self-repetition, and  $n$ -gram Vendi are independent evidence in a way the embedding metrics are not — and we win those too, including against the human-written bank on two of the four. The honest summary is therefore:

- **Embedding-space wins (Vendi, NN distance):** real but partly circular. Read them as confirmation that the optimizer works, not as independent evidence that the corpus is better.
- **Literal-space wins (duplication, distinct-2, self-repetition,  $n$ -gram Vendi):** independent, because nothing in the method targets them. These carry the argument.
- **Rendered-artifact measurements (CLIP, CLAP, MERT):** the most independent of all, since they live downstream of a second generative model the method never sees. This is why §7.6 matters more than its size in the paper suggests.

A reader who trusts only the third category still has the  $r = 0.155$  result, which is a finding about every method in the table rather than a comparison between them.

---

**I.7.11 Literal-space repulsion: fixing what the embeddings cannot see** Embedding metrics miss an entire class of repetition: tiled grids of one cell, prominent typography, a single shared palette, the same composition recolored. Audited with deliberately dumb, non-semantic signatures (a 16×16 luminance layout map, an autocorrelation tiling score, a hue histogram): half the images had a layout twin above 0.5 cosine, 40–45% were literal tilings, and 63–78% shared a palette — while prompt-level Jaccard sat at a healthy 0.14–0.17. Varied words, one visual mode: the conditional-dimension gap operating inside the *renderer*.

The fix is a four-quadrant repulsion —  $\{\text{text, vision}\} \times \{\text{literal, latent}\}$  — where the two literal quadrants were previously unpopulated: content-word overlap penalties and live overused-word bans on the text side; and on the vision side, measured structural bans injected into prompts (grids banned when recent

renders tile, dominant hue pairs named and banned, layout-change demands when layouts collide) plus a learned text  $\square$  bad-structure bridge that penalizes candidates near prompts whose renders tiled. Rerun at matched budget, palette twins halved (0.78  $\square$  0.35 in the coverage arm), the coverage objective improved 21% (0.206  $\square$  0.247), and layout/tiling moved modestly (0.51  $\square$  0.46–0.48; 0.45  $\square$  0.33) — a real but partial victory whose residue is the image model's own prior resisting text-side instruction, and which we report as such.

**I.7.12 Audio: the steered corpus, and embedder dependence** **The steered corpus.** 100 Lyria-3-Pro instrumentals generated by cross-modal steering with zero rejection — selection happens over prompts, every render is kept. Final measurements: 100% of tracks verify as instrumental in CLAP space (minimum instrumental-vs-vocal margin 0.077); mean-centered CLAP Vendi 17.43, against 11.5 (naive prompts) and 14.1 (axis-conditioned prompts) for the unsteered 47-track arms under the same embedder; opening loudness at 0.72 of each track's own median over the first three seconds, against a 0.46–0.61 baseline — the sparse-opening attractor substantially, not fully, suppressed.

**Structural control.** Lyria honors its own section-marker format ([ [A0]] [ [B1]] ...) and ignores wall-clock timestamps. Compliance with a requested section sequence is length-dependent: exact at prompt lengths near 560 characters, 0.11 at a mean of 658, zero by  $\sim$ 1,600. Structural contracts for music must be short or they are noise.

**Pre-render steerability is a property of the embedder, not of music.** Prompt-to-track alignment (Spearman between prompt-pair similarity through the text tower and track-pair similarity through the audio tower, naive/conditioned arms):

| embedder                            | alignment          | note                                                               |
|-------------------------------------|--------------------|--------------------------------------------------------------------|
| <b>MuQ-MuLan</b> (open MuLan-style) | <b>0.68 / 0.47</b> | a usable steering gradient                                         |
| CLAP-fused                          | 0.50 / 0.21        | weak; $\sim$ 0.05 on the steered corpus                            |
| CLAP-music                          | 0.18 / 0.06        | worst; median within-corpus NN distance 0.004 — it hears one track |
| MERT                                | —                  | no text tower (control)                                            |

Cross-embedder agreement on pairwise track similarity spans 0.22–0.84; per-arm diversity verdicts flip between embedders, and each embedder nominates a different most-redundant pair. Two prescriptions follow: steer music with MuQ-MuLan rather than CLAP, and never publish an audio-diversity number without naming its embedder.

## I.8 Relation to the coverage problem

Two objectives share the generation budget. **Max–min** is a packing objective:  $k$ -center-like, driven by the worst-case nearest pair, with no approximation guarantee under greedy selection. **Coverage** — given a finite budget, maximize the volume of a union of  $\epsilon$ -balls — is a covering objective, closer to facility location, and monotone submodular, so greedy selection carries a  $(1 - 1/e)$  guarantee. The sections above develop the max–min side; Part II develops the coverage side across two further domains.

Each objective's strength is the other's failure mode. Packing spreads points toward the boundary and over-invests in outliers, which is why max–min is vulnerable to the Theorem 4 pathology: outliers are what it is designed to seek. Covering weights regions by measure and fills the bulk, which is why it leaves the frontier sparse. Neither reduces to the other, and the tools that serve one degrade the other.

Both problems share this paper's structural constraints, and we expect the conditional-dimension results to be *more* consequential for coverage than for packing: covered volume is governed by reachable dimension, not by how many distinct prompts one can write, so the lattice-size / reachable-dimension gap of §6.1 bites harder there. The no-inverse-oracle assumption also degrades coverage's greedy guarantee, since the

guarantee is relative to the best subset *of what was proposed*, and the proposal distribution is confined to a slice.

---

## I.9 Limitations

**The simulations encode our hypothesis.** Mixture-of-modes with skewed weights and a rare diffuse junk component is a model of generator behavior, not a measurement of it. The theorems are unconditional given their assumptions and verified numerically; the *simulation* results inherit the model's assumptions. The live pilots are too small to independently confirm the asymptotics.

**One embedding oracle.** All geometry is nomic-embed-text geometry in the live runs. Whether “diverse” under one embedder transfers to another is untested here and is a real threat to any embedding-based diversity claim, including ours.

**Judge noise is modeled as unbiased.** Real LLM judges have systematic preferences — often *for* fluent, typical text, which is exactly the bias that would work against a diversity system. We model judge error as zero-mean Gaussian, which is optimistic.

**Additive spec composition.** §6.1 assumes conditioning attributes compose additively in embedding space. Real interactions between prompt attributes are not additive.

**The exam domain is simulated.** No real item bank was constructed, no psychometrician reviewed the live items, and  $\delta = 0.20$  as an enemy-item radius is a modeling choice, not a validated threshold.

**Quality is scalar.** Craft is not one number, and collapsing it to one lets a system trade away dimensions of quality the judge does not score.

---

## I.10 Conclusion

The infinite-horizon diversity objective is easy to write, cheap to maintain, and largely beside the point. What determines whether a corpus can keep growing without collapsing into paraphrase is the dimension of the generator's conditional support relative to the manifold it lives in, and every practically important consequence follows from that one gap: saturation arrives at rate  $n^{-1/m}$  rather than  $n^{-1/d}$ ; a single prompt covers an  $\varepsilon^{d-m}$  fraction of what the model could write; only transverse prompt motion raises the ceiling; and novelty pressure without a typicality constraint reliably escapes the manifold rather than exploring it.

Given an embedding oracle and no inverse, the system cannot compute its way to the next item. It can only choose what to condition on. The axis-scoring rule — spread, transversality, independence, headroom — is our answer to that choice, and its most valuable output is not the ranking but the **refine** signal: the moment it reports that nothing available is transverse any more is the moment the horizon has actually been reached, and the only remaining move is to ask the generator to subdivide its own vocabulary of variation.

That recursion is the single mechanism we found that changes the asymptote rather than the constant. Everything else — better selection, more candidates, higher temperature — buys a constant factor against a problem that is asymptotic, and the loudest of them buys it by quietly breaking the generator.

The measurements make the stakes concrete in a way the theory could not. A strong model, asked a perfectly reasonable question ten thousand times, returns the same item 1,637 times; the temperature knob barely moves that number and conditioning nearly erases it. Two diversity metrics computed on the same growing corpus point in opposite directions. And the text embeddings every method in this literature optimizes turn out to explain about two percent of the variance in whether the rendered images look alike. Each of those is a reason to distrust a single diversity number, and together they are the argument for the practice we ended up recommending: measure at the level of the artifact you are shipping, report the literal and the latent separately, publish the similarity distribution your kernel is operating on, and count your duplicates before you compute anything else.

---

---

## Part II — Coverage: reach as much as possible

### Abstract

Part I asks how to generate an unbounded stream of items that never collapses onto its own past, under a packing objective: keep every new item far from everything already made. This part studies the covering objective. A fixed budget of  $n = 10,000$  items must be spent so that the resulting set *covers* as much of the generator's reachable behavior space as possible: we maximize the measure of the union of small  $\epsilon$ -balls centered at the chosen embeddings, estimated by Monte Carlo against a reference pool of cheap draws from the generator itself. The coverage functional is monotone submodular, so greedy marginal-gain selection over a candidate stream inherits the classical  $(1 - 1/e)$  guarantee, and its Monte-Carlo estimate concentrates at the Hoeffding rate; we verify all three facts numerically rather than taking them on faith (0 violations of diminishing returns in 5,000 randomized chain trials; greedy attains the exact optimum on all 10 enumerable instances; the empirical 95% error of the pool estimate sits inside the Hoeffding band at every pool size tested). On two simulated domains at  $n = 10,000$  — red-team evaluation suites and user-persona corpora for dialogue testing — coverage-greedy selection covers 0.545 and 0.544 of a held-out 100k reachable pool at the calibrated mid radius, against 0.508 and 0.511 for naive iid and 0.388 and 0.367 for max-min packing, while max-min wins the min-gap metric by 20–35% and pays for it with a 3.2% / 2.9% junk rate against 0.0% for the gated policies: covering and packing are different objectives and optimizing either one visibly sacrifices the other. Recursive axis refinement roughly doubles coverage of the full latent space (0.381 vs 0.214; 0.538 vs 0.272) while *reducing* coverage of the fixed pre-refinement pool, which is the paper's central accounting problem: expanding the reachable space grows the denominator of the coverage ratio, so one run admits three defensible coverage numbers and we argue all of them must be reported. We also measure a *reachability gap* — the 17.4 points of coverage that no selection rule can recover because the proposal distribution cannot reach them, 92.5% of which recursive refinement does recover — which follows from having an embedding oracle but no inverse. The estimator is validated against analytically known answers, and evaluation is budget-matched on real gpt-5.6-luna corpora with a stratum-balanced reference and a scoring half held out from the selector. Coverage-greedy selection beats all five literature baselines -- SemDeDup, farthest-point/k-center, Self-Instruct's ROUGE-L filter, greedy MAP-DPP, and uniform random -- by 1.043x on DALL-E instructions and 1.091x on psychometric items. Against the released corpora (Alpaca, PersonaHub, WizardLM), covered fraction at a fixed radius is not comparable across corpora with different intrinsic scales -- it correlates -0.991 with within-corpus spacing and -0.986 with Vendi -- so we use a scale-free estimator with reference-side k-NN radii. On it, coverage-retrieval conditioning places first of twelve corpora at 0.4441 against Alpaca's 0.3722, at one-twentieth of Alpaca's generation budget and with the highest precision in the field (0.973). The same data shows the covering/packing dissociation at full strength: k-center maximizes min-gap (0.696) and finishes last on coverage (0.096, a fifth of random), and the two highest-Vendi selectors are the two worst covering ones. Naive psychometric prompting is 73.6% exact duplicates (one item repeated 2,726 times in 10,000 generations), which high temperature barely dents (71.0%) and conditioning nearly eliminates (0.0–4.7%) -- direct evidence that conditioning, not temperature, buys coverage, and a demonstration that nearest-neighbor epsilon-calibration is undefined on such a corpus (median self-distance exactly 0). Finally, literal and latent diversity move in opposite directions as the corpus grows (distinct-2 falls  $0.944 \square 0.770$  while mean-centered Vendi rises  $3.98 \square 37.06$ ), and the uncentered Vendi is compressed 2.7 $\times$  by the embeddings' shared mean direction — so both levels, and both centerings, need reporting.

### II.1 Introduction

A team building an evaluation suite for a language model does not get an infinite horizon. They get a budget — ten thousand prompts, say — and a product requirement that is naturally a *coverage* requirement: the suite should contain an example near every qualitatively distinct behavior the system under test might

encounter, because a suite that clusters into a few families gives false assurance about everything outside them. The same shape of problem appears when generating synthetic user personas to exercise a dialogue system: near-duplicate personas waste budget, and bugs live in the parts of user-behavior space the corpus never visits.

The parent paper in `research/infinite_horizon_diversity/` (this work lives in its `coverage/` subdirectory and imports its modules directly) studies the open-ended version of this problem: an unbounded stream, an objective that combines typicality (mean distance to prior embeddings) with max-min novelty, and a method stack — generator-elicited latent axes, spectral orthogonalization against the corpus second-moment matrix, recursive refinement of saturated cells, and an append-only attractor ledger — that keeps the stream from collapsing at any horizon. This paper keeps the method stack and swaps the objective, because the budgeted problem is *not* the truncation of the infinite one:

1. **Max-min packing is the wrong objective at a finite budget.** Packing pushes points apart; its optima sit on the boundary of the reachable region and over-invest in extremes and outliers, because an outlier is by construction far from everything. Coverage weights regions by how much probability mass they hold and fills the bulk first. The two objectives correspond to two different classical problems —  $k$ -center versus maximum coverage / facility location — and their optima genuinely differ (§3.4, §5).
2. **A known  $n$  permits planning.** With an unbounded horizon there is nothing to allocate; with  $n = 10,000$  known in advance one can ask how the budget should be split across cells of the behavior space — equally, proportionally to measure, or adaptively — and the answer depends on the radius  $\epsilon$  at which coverage is demanded (§3.5, §6).
3. **Coverage needs a denominator.** “We covered 62% ” is meaningless without saying 62% of *what*. We measure coverage against the generator’s own reachable distribution, approximated by a pool of cheap draws — and this choice has a sting: Part I’s recursive refinement *expands* the reachable space, so the better the method is at opening new territory, the larger the space it is graded against (§7).

Everything is built to the same scaling discipline as Part I: per-item cost is  $O(1)$  in  $n$  (fixed-size reference pool, incremental covered masks, bounded history subsamples), so nothing in the pipeline gets slower at item 10,000 than at item 100.

We make five contributions. First, a formalization of budgeted corpus construction as Monte-Carlo maximization of a union-of-balls coverage functional, with the relevant guarantees stated precisely and *checked numerically* (§3, §6), including the no-inverse-oracle propositions and a direct measurement of the reachability gap they imply (§2, §6.6). Second, a simulation study on two ground-truth mixture worlds — red-team suites and persona corpora — comparing naive iid sampling, max-min packing, coverage greedy, gated coverage, and gated coverage with recursive refinement at  $n = 10,000$  (§5). Third, a coverage adaptation of Part I’s axis-scoring rule in which the four-factor axis-scoring product collapses to a single estimable scalar — expected marginal  $\epsilon$ -ball gain of conditioning — with an ablation isolating what score-guided conditioning buys (§4, §5.3). Fourth, a live pilot of the full method (elicited axes, score-guided focus, judge gate, ledger, refinement, coverage-greedy selection) on gpt-5.6-luna with local embeddings (§8). Fifth, an explicit treatment of the honest-denominator problem that refinement creates (§7).

## II.2 Problem statement

Fix an embedding map  $\phi$  into  $\mathbb{R}^D$  (we use unit-normalized embeddings,  $D = 768$  live,  $D = 32$  in simulation) and a radius  $\epsilon > 0$ . The generator, prompted in whatever ways the pipeline can reach, induces a *reachable distribution*  $\mu$  over  $\mathbb{R}^D$ : the pushforward of everything the generator can actually be made to produce. Given a budget  $n$ , choose a set  $S = \{x_1, \dots, x_n\}$  of generated items to maximize the covered measure

$$F_{\mu, \epsilon}(S) = \mu(\bigcup_{x \in S} B(x, \epsilon)),$$

the probability that a fresh draw from the generator lands within  $\epsilon$  of some chosen item. Equivalently,  $S$  should be an  $\epsilon$ -net for as much of  $\mu$  as the budget allows. Three deliberate choices:

- **Coverage of the reachable space, not of  $\mathbb{R}^D$ .** Balls in empty space cover Lebesgue volume but no behavior. Measuring against  $\mu$  makes the metric mean “fraction of what the generator can do that the

corpus has an exemplar of” — which is the product requirement for an eval suite. The cost of this choice is the denominator problem of §7.

- **$\epsilon$  is a resolution parameter, not a nuisance.** Small  $\epsilon$  asks for exemplars of fine behavioral distinctions; large  $\epsilon$  only for one exemplar per coarse region. Every result below is reported at three radii calibrated per domain (§5.1), because the ranking of methods changes with  $\epsilon$ .
- **Quality is a constraint, not part of the objective.** An item that covers new territory but fails a quality / typicality bar is not an asset; §5.4 quantifies what happens when the constraint is dropped.

**The oracle asymmetry.** One assumption structures everything downstream: we have an embedding oracle  $E: \text{text} \rightarrow \mathbb{R}^D$ , and *no inverse*. There is no  $E^{-1}$  that decodes a target point back into text. We can compute exactly where in embedding space the next item should land — the largest uncovered region, the trailing eigenspace — but we cannot manufacture an item that lands there. The only handles a pipeline has are (a) language-valued latent conditioning that *steers* the generator's proposal distribution, and (b) selection among sampled candidates. This is why the architecture is propose–measure–select rather than solve, and it has two consequences we state as propositions:

**Proposition 1 (no direct construction).** Coverage-optimal point sets are not directly constructible. The greedy  $(1 - 1/e)$  guarantee (§3.2) holds relative to the best subset of the *proposal distribution's support*, not of  $\mathbb{R}^D$ ; the difference between what an unconstrained selector could cover and what a proposal-limited selector can is a *reachability gap*, which is a property of the generator-plus-conditioning stack, not of the selection algorithm. We measure it directly in §6.6.

**Proposition 2 (conditioning as surrogate inverse).** Language-valued conditioning is the only available surrogate for  $E^{-1}$ : choosing which axis to condition on next is choosing which *slice* of the reachable manifold the next candidates will be sampled from, so a axis scoring that ranks axes by where their slices land (§4) is precisely an approximate inverse — it maps “where we want mass” to “which words to condition on”.

## II.3 Theory

**II.3.1 The coverage functional is monotone submodular** For any measure  $\mu$  and radius  $\epsilon$ , define  $F(S) = \mu(\bigcup_{x \in S} B(x, \epsilon))$  on finite  $S \subseteq \mathbb{R}^D$ .  $F$  is monotone (adding a ball cannot uncover anything) and submodular: for  $S \subseteq T$  and any  $x$ ,

$$F(S \cup \{x\}) - F(S) \geq F(T \cup \{x\}) - F(T).$$

*Proof sketch.* The marginal gain of  $x$  given  $S$  is  $\mu(B(x, \epsilon) \setminus \bigcup_{y \in S} B(y, \epsilon))$ . Since  $S \subseteq T$  implies  $\bigcup_{y \in S} B(y, \epsilon) \subseteq \bigcup_{y \in T} B(y, \epsilon)$ , the set being subtracted only grows, so the gain only shrinks. This is the standard coverage-function argument; nothing about balls is essential — any map from items to measurable “footprints” yields a submodular  $F$ .  $\square$

The same argument applies verbatim to the Monte-Carlo estimator we actually optimize. Given a reference pool  $P = \{p_1, \dots, p_P\}$  of iid draws from  $\mu$ ,

$$\hat{F}(S) = (1/P) \cdot \#\{p \in P : \text{dist}(p, S) \leq \epsilon\}$$

is itself a finite coverage function — item  $x$  covers the fixed subset  $N_\epsilon(x) = \{p : \|p - x\| \leq \epsilon\}$  of the pool, and  $\hat{F}(S) = |\bigcup_{x \in S} N_\epsilon(x)| / P$  — so  $\hat{F}$  is monotone submodular *exactly*, not merely in expectation. Our implementation check (§6.1) found 0 violations of either property in 5,000 randomized nested-chain trials, as it must if the code is right.

**II.3.2 Greedy and its guarantee** Because  $\hat{F}$  is monotone submodular with  $\hat{F}(\emptyset) = 0$ , the greedy algorithm that repeatedly adds the item of maximum marginal gain satisfies the Nemhauser–Wolsey–Fisher bound

$$\hat{F}(S_{\text{greedy}}) \geq (1 - 1/e) \cdot \max_{|S| \leq n} \hat{F}(S) \approx 0.632 \cdot \text{OPT},$$

and this is tight for the class (maximizing coverage is NP-hard, and beating  $1 - 1/e$  is hard under standard assumptions, Feige 1998). §6.3 checks the bound against *exact* optima on instances small enough to enu-

merate; greedy attained the optimum itself on all ten instances, comfortably above the bound — typical behavior, since the  $1 - 1/e$  worst case requires adversarial structure.

**II.3.3 Sequential generation is streaming greedy** A generation pipeline cannot pick from a precomputed ground set: candidates arrive a few at a time ( $K$  per step) and must be accepted or discarded immediately. This is the streaming submodular maximization regime. Sieve-streaming (Badanidiyuru, Mirzasoleiman, Karbasi, Krause 2014) maintains geometric thresholds and accepts any arrival whose marginal gain clears one, achieving a  $(1/2 - \delta)$  guarantee in one pass with  $O((k \log k)/\delta)$  memory, independent of stream length. Our per-step rule — take the best of  $K$  fresh candidates by marginal gain — is the natural best-of- $K$  relaxation: weaker in the worst case (an adversarial stream can starve it) but stronger on exchangeable streams like ours, where each batch is iid from the current proposal distribution and thresholding would only discard budget. Empirically the  $K = 4$  stream rule retains most of full greedy's value over an identical ground set (0.484 vs 0.528 covered fraction; §6.2) at a per-item cost that never touches the whole ground set. What matters for scaling is that the marginal-gain oracle is  $O(K \cdot P)$  per step against a fixed-size pool with an incrementally maintained covered mask —  $O(1)$  in  $n$ .

**II.3.4 Coverage is not packing** Max-min selection (choose the candidate farthest from the accepted set) is the greedy algorithm for the  $k$ -center / packing objective, and the sibling paper shows it is the right *shape* of objective for an unbounded stream once a typicality anchor keeps it on-manifold. At a finite budget against a fixed  $\mu$  the two objectives pull apart:

- $k$ -center cares about the worst-covered *point of  $S$  itself*; its optima push to the boundary of the support and place points in regions of vanishing measure (an outlier is far from everything, so packing *rewards* it — the novelty-junk conflation of Part I, now with a budget price attached: every junk item is a ball spent covering  $\approx 0$  measure).
- Maximum coverage / facility-location-type objectives weight territory by measure: a second exemplar in a heavy mode can be worth more than a first exemplar of a negligible one, and junk is worth nothing because  $\mu$  puts almost nothing near it.

§6.2 isolates the contrast with no generator in the loop: on one shared candidate set, greedy-coverage covered 0.528 of a held-out pool where farthest-point packing covered 0.143 — while packing's min-gap (1.44) beat greedy's (0.577) by 2.5 $\times$ . Neither is "better"; they optimize different functionals, and §5 shows the same double dissociation with the full generation loop in place.

**II.3.5 Estimation error of the Monte-Carlo coverage** For fixed  $S$ , each pool point contributes an iid Bernoulli( $F(S)$ ) indicator, so Hoeffding / Chernoff gives

$$P(|\hat{F}(S) - F(S)| > t) \leq 2 \exp(-2 P t^2),$$

i.e. a 95% band of  $t_{95} = \sqrt{\ln(2/0.05)/(2P)}$  — about  $\pm 0.030$  at  $P = 2,000$  and  $\pm 0.015$  at  $P = 8,000$ . Two caveats we take seriously. First, the bound is for fixed  $S$ ; a set *selected* by optimizing  $\hat{F}$  on the same pool is biased upward on that pool (the winner's curse), which is why every simulation number in §5 is reported on held-out pools the selector never saw, and the live pilot reports selection-pool coverage explicitly labeled as such. Second, uniform-over- $S$  guarantees would need a union bound over the (exponentially many) candidate sets; we do not rely on one — held-out evaluation makes the fixed- $S$  bound the relevant one. §6.4 confirms the band empirically: at every pool size tested,  $\geq 99\%$  of 200 independent pool estimates fell within  $t_{95}$  of a 200k-pool ground truth.

**II.3.6 Allocation under a known budget** Suppose the space decomposes into cells (modes) with measures  $w_m$ , and within-cell placement is near-saturating: with  $c_m$  items in cell  $m$ , the covered measure of that cell is approximately  $w_m \cdot g_m(c_m)$  with  $g_m$  increasing and concave (each additional exemplar covers a shrinking uncovered remainder — submodularity again, now per cell). Maximizing  $\sum w_m g_m(c_m)$  subject to  $\sum c_m = n$  is a concave allocation problem whose optimum is the water-filling rule: equalize marginal gains  $w_m g'_m(c_m)$  across cells. Three closed-form rules bracket the regimes:

- $\epsilon$  **large relative to cell diameter** ( $g_m$  saturates after  $\sim 1$  item): any rule that puts at least one item per cell is near-optimal; equal allocation wastes least; proportional over-invests in heavy cells.
- $\epsilon$  **small** ( $g_m$  near-linear over feasible  $c_m$ , slope  $\square$  per-item covered mass  $\approx$  density  $\times$  ball volume): marginal gain is  $w_m$ -weighted throughout, and proportional-to-measure allocation dominates equal.
- **In between**, neither closed form matches water-filling; greedy marginal coverage *is* water-filling implemented empirically, and wins at the radius it optimizes.

The measured allocation experiment (§6.5) shows exactly this pattern, plus a failure mode the clean theory hides: once greedy has saturated the pool at its selection radius, its marginal signal is identically zero and the tie-breaking rule silently decides where the rest of the budget goes.

**Depth per condition and switch cost.** Part I's slice theory sharpens the within-cell side of allocation. For a *fixed* prompt/spec the generator concentrates on an  $m$ -dimensional slice with  $m \square d$ , so fixed-prompt novelty decays as  $n^{(-1/m)}$  (fitted slopes  $-0.472, -0.309, -0.195$  for  $m = 2, 3, 5$  — governed by  $m$ , not by the ambient dimension), and one slice  $\epsilon$ -covers a vanishing  $\epsilon^{(d-m)}$  fraction of the manifold. The budgeted consequence is a breadth-vs-depth rule whose correct form (their P3, corrected from its first statement) is: with *free* condition-switching the optimum is the boundary  $n^* = 1$  — one sample per spec, then move — and an interior optimal depth exists only once switching carries a real cost  $c$  (eliciting a spec, embedding levels, an occasional refinement call): their measured  $n^*$  rises from 1 ( $c = 0$ ) to 10 ( $c = 3$ ) to 30 ( $c = 30$ ). Optimal depth is set by the switch-cost-to-sample-cost ratio, not by  $\epsilon$  alone. Our pipelines sit near the cheap-switching regime — in simulation a spec switch is free, and in the live pilot we measure  $c \approx 0.10$  generation-equivalents (§8: 5 axis-elicitation, refinement and ledger-mining calls amortized over 50 steps, against 3 generations per step) — which is why every policy here draws its  $K$  candidates from  $K$  *fresh* specs rather than sampling any spec deeply. At that switch cost the sibling's sweep puts the optimum at or adjacent to  $n^* = 1$ , which is what we do.

## II.4 Method

The pipeline is Part I's stack with the selection objective and its bookkeeping swapped from packing to covering. Concretely:

**Reachable pool (the denominator).** Before selection begins, draw a pool of cheap, unconditioned samples from the generator and embed them. This pool *is* the Monte-Carlo measure: coverage means coverage of it. In simulation the pool has 20,000 points (selection) plus held-out pools of 20,000 (tracking) and 100,000 (final evaluation); live, 240 one-shot prompts. Cost is  $O(P)$  once, amortized over the whole run.

**Elicited axes and specs.** The generator is asked — before generating any items — to name the latent axes along which items in this domain can differ, each with discrete language-valued levels (live: 6 axes for evaluation prompts, e.g. persona pressure, trap construction, response obligation). A spec is one level per applicable axis;  $K$  specs are sampled per step and one candidate is generated under each. Axes give the proposal distribution spread the base generator lacks; they are the reason the method's candidates are worth ranking at all.

**Score-guided conditioning (the coverage adaptation).** Part I's generation axis scoring ranks candidate axes by a four-factor product — spread (are the axis's levels different from each other), transversality (does their variation point outside the corpus's occupied eigenspace), independence (measured against axes already in use, never against rival candidates — scoring candidates against each other makes duplicates zero each other out), and headroom (an entropy deficit over the axis's sampled levels). Under a *coverage* objective this product collapses, and the collapse is itself a packing-vs-covering result. Packing has no per-axis marginal decomposition — min-gap is a global property — so the axis scoring must approximate “will this axis move the slice somewhere new” from four separate geometric proxies. Submodular coverage has an exact marginal quantity: the **expected marginal  $\epsilon$ -ball gain of conditioning on the axis**,  $E_{\{x \square \text{axis}\}}[F(S \cup \{x\}) - F(S)]$ . That one scalar subsumes the four factors: an axis with near-synonym levels (no spread) or whose slices sit inside covered territory (no transversality) or which duplicates an axis already exploited (no independence, because the covered mask already contains that axis's contribution and re-scores every candidate after every accept — the same greedy re-scoring insight as the axis scoring's `select_axis_set`, executed by the objective itself) or whose region is already covered (no headroom) all have small expected marginal gain,

and the headroom that remains is *measure-weighted* by construction — a large under-sampled region beats a small one because it holds more uncovered pool mass. In simulation we estimate the scalar directly: each spec carries an estimate initialized from an 8-draw probe at creation and updated as an EMA of observed candidate gains, and the K specs per step are drawn proportionally to it with a 20% uniform exploration floor (`gated_coverage_refine_calc`). In the live pilot, where axes rather than opaque specs are the unit, we keep the sibling’s embedding-based spread and transversality (level-description embeddings against the accepted corpus’s occupied eigenspace) and replace its entropy headroom with the measure-weighted estimate (per-level mean observed marginal gain, optimistic for unused levels); the top-two axes by this promise become the step’s dominant design pressures, with their most-different levels chosen by farthest-point selection in level-embedding space — the axis scoring acting as the surrogate inverse of Proposition 2.

**Coverage-greedy selection.** Each step: embed the K candidates, compute each one’s marginal gain — the number of not-yet-covered pool points within  $\epsilon_{\text{sel}}$  — and accept the gated candidate of maximum gain (ties broken by judge score). The covered mask is updated incrementally; per-step cost is  $O(K \cdot P + P)$  regardless of n.

**Judge gate (quality/typicality).** A separate LLM call (generation never self-grades) scores each candidate on a rubric; candidates below the bar are ineligible regardless of gain. In simulation the gate also enforces a typicality z-score against the running centroid, mirroring the sibling paper’s anchor. The gate is negative supervision only: if every candidate fails, the step falls back to best-judged rather than emitting nothing, and the rejection is logged.

**Attractor ledger.** Every 20 accepted items (live; 50 in the sibling paper), the accepted corpus is shown to the model with the question “what do these have in common?”; the mined attractors are appended to a JSONL ledger, and a bounded top-K slice becomes explicit negative constraints in subsequent generation prompts. The ledger is append-only; its influence on any single prompt is fixed-size.

**Recursive refinement + pool refresh.** When recent accepted items’ marginal gains saturate (the region’s balls mostly re-cover covered ground), the responsible region of spec-space is described back to the generator, which either splits an axis level into finer sub-levels or mints a new axis meaningful inside that cell — conditional axes forming a tree, exactly as in Part I. The budgeted setting adds one obligation the infinite-horizon setting does not have: refinement changes the reachable distribution, so the reference pool must be refreshed with draws from the newly opened region, or the selector will see zero gain precisely where the new territory is (its balls cover no *old* pool points). This pool-refresh step is the operational face of the honest-denominator problem (§7).

## II.5 Simulation study

**II.5.1 Worlds** Both domains instantiate Part I’s MixtureWorld pattern with a spec layer added between policy and world, so that policies only ever touch the handles a real pipeline has. A world has M latent modes (mixture of anisotropic Gaussians, centers at norm 1.5, per-mode  $\sigma \in [0.10, 0.22]$ ) with skewed Dirichlet(0.6) weights, latent per-mode quality in [0.5, 0.9], a noisy judge ( $\sigma = 0.10$ ) over it, and a rare diffuse junk component (weight 0.8%,  $\sigma = 3.0$ , quality  $\approx 0.05$ ) representing off-manifold text. The policy addresses the generator through *specs*: opaque handles that map (invisibly) to modes. Crucially, base specs reach only a subset of modes — 25 of 60 (redteam), 20 of 40 (persona) — because a real generator’s unconditioned repertoire does not span its latent space; refinement is the only way to grow reach.

- **redteam** (evaluation-suite world): M = 60 attack families, full-dimensional scatter in D = 32.
- **persona** (dialogue-testing world): M = 40 archetypes whose centers lie on a random 8-dimensional linear patch of  $R^{32}$  (low intrinsic dimension, as embedded persona text exhibits), with small ambient jitter.

Radii are self-calibrated per domain:  $\delta$  = median nearest-neighbor distance from held-out reachable draws to n = 10,000 iid reachable draws, and results are reported at  $\epsilon \in \{0.6\delta, \delta, 1.6\delta\}$ . By construction naive iid covers  $\approx 50\%$  of the pool at  $\epsilon = \delta$ , which makes the mid radius interpretable on sight. Measured:  $\delta = 0.887$  (redteam), 0.787 (persona).

**II.5.2 Policies** All policies run at  $n = 10,000$  with  $K = 4$  candidates per step (naive takes the first draw), identical world geometry, and disjoint RNG streams:

| policy                     | selection rule                                               | gate                                       | n |
|----------------------------|--------------------------------------------------------------|--------------------------------------------|---|
| naive                      | first draw                                                   | –                                          | – |
| maxmin                     | argmax min-dist to accepted set (bounded 1,000-subsample)    | –                                          | – |
| coverage                   | argmax marginal $\epsilon$ -ball gain on 20k estimation pool | –                                          | – |
| gated_coverage             | same                                                         | judge $\geq 0.35 \wedge$ anchor $z \leq 3$ | – |
| gated_coverage_refine      | same                                                         | same                                       | s |
| gated_coverage_refine_calc | same                                                         | same                                       | s |

**II.5.3 Results** Covered fraction of the held-out 100k base-reachable pool at the calibrated mid radius (redteam  $\epsilon = 0.887$  / persona  $\epsilon = 0.787$ ), with the coverage of the full latent space (the oracle pool base policies cannot reach), junk, quality, and the packing metric, all at  $n = 10,000$ :

**redteam** ( $\delta = 0.887$ ):

| policy                            | covered base @ $\epsilon_{\text{mid}}$ | covered world @ $\epsilon_{\text{mid}}$ | junk %      | mean quality | min-gap      | Vend |
|-----------------------------------|----------------------------------------|-----------------------------------------|-------------|--------------|--------------|------|
| naive                             | 0.508                                  | 0.195                                   | 0.70        | 0.684        | 0.430        | 23.8 |
| maxmin                            | 0.388                                  | 0.159                                   | 3.19        | 0.690        | <b>0.549</b> | 23.8 |
| coverage                          | <b>0.545</b>                           | 0.214                                   | 0.70        | 0.681        | 0.458        | 23.4 |
| gated_coverage                    | 0.541                                  | 0.214                                   | <b>0.00</b> | 0.761        | 0.403        | 22.2 |
| gated_coverage_refine             | 0.492                                  | 0.381                                   | <b>0.00</b> | 0.765        | 0.056        | 18.8 |
| <b>gated_coverage_refine_calc</b> | <b>0.567</b>                           | <b>0.427</b>                            | <b>0.00</b> | 0.760        | 0.138        | 23.9 |

**persona** ( $\delta = 0.787$ ):

| policy                            | covered base @ $\epsilon_{\text{mid}}$ | covered world @ $\epsilon_{\text{mid}}$ | junk %      | mean quality | min-gap      | Vend |
|-----------------------------------|----------------------------------------|-----------------------------------------|-------------|--------------|--------------|------|
| naive                             | 0.511                                  | 0.257                                   | 0.90        | 0.707        | 0.431        | 14.5 |
| maxmin                            | 0.367                                  | 0.191                                   | 2.91        | 0.705        | <b>0.549</b> | 16.3 |
| coverage                          | <b>0.544</b>                           | 0.276                                   | 0.51        | 0.702        | 0.443        | 14.0 |
| gated_coverage                    | 0.537                                  | 0.272                                   | <b>0.00</b> | 0.786        | 0.410        | 13.1 |
| gated_coverage_refine             | 0.501                                  | 0.538                                   | <b>0.00</b> | 0.787        | 0.064        | 9.0  |
| <b>gated_coverage_refine_calc</b> | <b>0.571</b>                           | <b>0.582</b>                            | <b>0.00</b> | 0.757        | 0.102        | 11.5 |

Four observations, consistent across both domains.

1. **The double dissociation is real with the full generation loop in place.** Max-min wins its own objective decisively (min-gap 0.549 vs  $\approx 0.43$ – $0.46$  for everything else) and pays 12–18 points of covered fraction for it — *below naive iid* on the coverage objective in both domains, because its budget migrates to boundaries and outliers where the reachable measure is thin. Coverage-greedy beats naive by 3.3–3.7 points at  $\epsilon_{\text{mid}}$  on the held-out pool; the gap versus maxmin is 15.7 (redteam) and 17.6 (persona) points.
2. **Radii change the ranking.** At the large radius ( $1.6\delta$ ) every policy is  $\approx 0.98$ – $0.99$  on the base pool — the base reachable set is easy to cover coarsely, and differences live at and below  $\delta$ . At the small radius ( $0.6\delta$ ), refinement dominates in redteam (0.031 vs 0.014 naive: tight child specs supply fine-grained exemplars) while in persona the low-intrinsic-dimension geometry makes small balls so poor at catching measure (covered fractions of  $10^{-3}$ ) that no policy separates meaningfully — resolution demands must respect the manifold dimension (see also §6.7).

3. **Refinement trades base-pool coverage for reach.** The refine policies give back  $\sim 4$  points on the fixed base pool (0.49–0.50 vs 0.54 gated) and in exchange nearly double coverage of the full latent space (0.381 vs 0.214 redteam; 0.538 vs 0.272 persona), reaching all 60/60 and 40/40 modes versus the 25/60 and 20/40 the base specs can address. Which of those numbers is “the” coverage of the run is exactly the denominator question of §7.
4. **Diversity metrics disagree, instructively.** The refine policies have the *lowest* Vendi scores (18.8 and 9.0) while covering the most latent territory — linear-kernel Vendi measures global spectral spread, and the dense local clusters that tight child specs produce (min-gap 0.056) concentrate the spectrum even as they open new regions. A single diversity number, whether min-gap or Vendi, cannot summarize a coverage run; the covered-fraction-by-denominator table can. This is the coverage-side face of Part I’s metric trap (§11): its fixed-prompt policy scored *best* on residual-headroom precisely because it covered nothing.
5. **Score-guided conditioning dominates, and repairs refinement’s regression.** Adding promise-weighted spec sampling (§4) to the refine policy is the only change that wins on *both* denominators at once: it is the best policy on the fixed base pool (0.567 / 0.571, ahead of plain coverage-greedy at 0.545 / 0.544) *and* the best on the full latent space (0.427 / 0.582, ahead of unguided refinement at 0.381 / 0.538). It does this while recovering most of what unguided refinement gave up elsewhere: Vendi rises from 18.8 to 23.9 (redteam) and 9.0 to 11.5 (persona), and min-gap from 0.056 to 0.138 and 0.064 to 0.102. The mechanism is straightforward — unguided refinement spends budget uniformly across child specs, most of which land in already-covered territory, whereas promise-weighted sampling concentrates draws on the specs whose expected marginal  $\epsilon$ -ball gain is still large. The ablation isolates this: the two policies differ *only* in how specs are sampled, with identical gates, identical refinement triggers, and identical pool refresh.

**II.5.4 Junk, quality, and the gate** The junk column isolates the novelty–junk conflation at a finite budget. Max-min accepts 3.2% / 2.9% junk — with  $K = 4$  candidates and junk weight 0.8%, essentially every junk draw that appears among the candidates is selected, because a diffuse off-manifold point maximizes distance to everything; each such acceptance is a ball spent covering  $\approx 0$  reachable measure. Ungated coverage does *not* share the attraction: junk lands at or below the naive base rate (0.70% / 0.51%), because a ball around junk covers almost no pool mass, so junk wins the argmax only when every candidate’s marginal gain is  $\approx 0$  (late-run ties). The gate removes junk entirely (0.00% over 10,000 accepts in all four gated runs) and raises mean true quality by 0.08 in both domains at a cost of  $\leq 0.7$  points of covered fraction — the cheap insurance Part I’s typicality gate promised, now priced under a budget. The residual case for the gate under coverage is therefore not junk *attraction* (packing’s failure) but junk *indifference*: a quality-blind coverage selector will still spend late-run ties on garbage, and mean quality drifts with whatever the proposal distribution emits.

## II.6 Numerical verification of the theory

All numbers in figures/summary\_theory.json.

**6.1 Submodularity (exact).** 200 candidates, 5,000-point pool,  $\epsilon = 1.0$ . Over 5,000 random nested chains  $S \sqsubseteq T$  with a fresh  $x$ : 0 violations of diminishing returns, 0 of monotonicity, worst violation 0.0.

**6.2 Selection rules head-to-head (no generator).** One shared ground set of 2,000 candidates, budget  $k = 300$ ,  $\epsilon = 1.0$ , held-out 20k pool:

| rule                   | covered fraction (held-out) | min-gap |
|------------------------|-----------------------------|---------|
| full greedy            | 0.528                       | 0.58    |
| stream greedy, $K = 4$ | 0.484                       | 0.68    |
| random                 | 0.420                       | 0.53    |
| max-min (Gonzalez)     | 0.143                       | 1.44    |

Packing is catastrophic *as a coverage algorithm* (below random by 3×) while being unbeatable at its own metric — the pure-form double dissociation. Stream greedy with  $K = 4$  keeps 92% of full greedy's coverage.

**6.3 The  $(1 - 1/e)$  bound against exact optima.** Ten instances with  $|\text{ground}| = 30$ ,  $k = 5$  (142,506 subsets enumerated each): greedy/OPT ratio was 1.0 on every instance; bound 0.632 never approached.

**6.4 Monte-Carlo error.** Fixed  $S$  of 300 items, ground truth from a 200,000-point pool ( $F = 0.4134$ ). Over 200 independent pools per size: empirical 95th-percentile absolute error 0.042 / 0.020 / 0.011 at  $P = 500$  / 2,000 / 8,000, versus Hoeffding  $t_{95} = 0.061$  / 0.030 / 0.015; the fraction of estimates inside the band was  $0.99\text{--}0.995 \geq 0.95$  everywhere.  $P = 20,000$  (our estimation pool) implies  $t_{95} \approx 0.0096$ : pool noise is well below every effect size we interpret.

**6.5 Allocation rules.**  $n = 10,000$  across the redteam world's  $M = 60$  cells with *known* measures (the oracle planning question), evaluated against a 50k full-mixture pool:

| rule                                   | $\epsilon = 0.6$ | $\epsilon = 1.0$ | $\epsilon = 1.6$ | $\epsilon = 2.4$ |
|----------------------------------------|------------------|------------------|------------------|------------------|
| equal per cell                         | 0.046            | 0.494            | 0.9916           | 0.9921           |
| □ measure                              | <b>0.060</b>     | 0.547            | 0.9917           | 0.9921           |
| □ sqrt(measure)                        | 0.055            | 0.531            | 0.9918           | 0.9921           |
| greedy marginal (at $\epsilon = 1.0$ ) | 0.004            | <b>0.618</b>     | 0.9919           | 0.9921           |

Three regimes, as §3.6 predicts. At large  $\epsilon$  every rule saturates (all  $\approx 0.992$  — the remaining 0.8% is junk measure no allocation can reach). At the greedy-optimized radius, greedy beats the best closed form by 7 points: water-filling in the wild. At small  $\epsilon$ , proportional wins among closed forms (near-linear per-cell returns favor measure-weighting) — and greedy at  $\epsilon = 1.0$  is *terrible* (0.004), for an instructive reason: greedy saturated its  $\epsilon = 1.0$  pool after roughly half the budget, its marginal signal went identically to zero, and argmax tie-breaking silently dumped 5,083 of 10,000 items into one arbitrary cell (correlation of greedy's allocation with every closed form  $\leq 0.17$ ). Lesson for practitioners: greedy's guarantee says nothing about what it does *after* it has won; at a finite budget you must either select at the smallest  $\epsilon$  you care about, or hand the post-saturation residual to an explicit rule (e.g. proportional), or lower  $\epsilon$  adaptively as gains vanish.

**6.6 The reachability gap (Proposition 1, measured).** Identical greedy selection ( $\epsilon = 1.0$ ,  $k = 300$ ,  $|\text{ground}| = 2,000$ ), three proposal distributions, one held-out full-mixture pool. An *oracle* ground set drawn from the full latent mixture — the best a selector could do if conditioning could reach everything — covers 0.387. A *proposal-limited* ground set drawn from the base specs (which reach 25 of 60 modes) covers 0.212: a reachability gap of 17.4 points that no selection algorithm can close, because the deficit is in what the generator can be made to emit, not in how candidates are chosen. A *refined-proposal* ground set (one spec per mode — the target recursive refinement works toward) covers 0.374, recovering 92.5% of the gap. This is Proposition 1 made quantitative: the  $(1 - 1/e)$  guarantee binds relative to the proposal support, and conditioning — not selection — is where the missing coverage lives.

**6.7 Lattice size is not reachable dimension.** Using Part I's slice-structured world, in which specs compose level directions additively, we build two worlds with *identical* lattice size  $L^A = 1,024$  distinct specs but different rank of the level-direction span: full-rank (rank 10, reachable dimension 13) versus low-rank (all axes forced into a shared 2-dim subspace: rank 2, reachable dimension 5). Same budget ( $n = 2,000$  random-spec draws), same radius ( $\epsilon = 0.27$ ): the low-rank world's own reachable pool is 68.1% covered where the full-rank world's is 49.0%. The number of distinct prompts you can write says nothing about how much there is to cover; covered volume is governed by reachable dimension, and a huge lattice over low-rank level directions is a thin space wearing a combinatorial costume. For coverage this distinction is sharper than for packing — the denominator itself is set by reachable dimension — and it is why our simulation worlds put a spec layer between policy and world rather than drawing from a global mixture: a simulator without conditioning structure cannot exhibit the effect and understates refinement (the sibling paper measured exactly this failure in its own first simulator).

**6.8 Refinement direction: density vs reach.** Part I's negative result — naive refinement hurt, because tight children near saturated parents concentrate mass where the corpus already sits — ports to coverage,

with a measurement subtlety worth recording. Nearest-neighbor distance CANNOT distinguish density-refinement from reach-refinement: in our first version of this check, child draws sat  $\approx 1.04$  from the nearest corpus point whether they had moved in-span or off-span, because a 1,500-item corpus is sparse in its own 13-dim occupied span. Density-vs-reach is a *span* property. Measured with span-level metrics: *local* children (tight,  $0.25\times$  displacement — the naive move) produce draws of which 36.7% are already covered by the pre-refinement corpus at the coverage radius, and add  $+0.02$  effective dimensions — almost pure density. Full-length *isotropic* displacement produces 0% pre-covered draws and  $+0.13$  effective dimensions; *transversality-guided* displacement 0% and  $+0.14$ , with off-span displacement energy 0.915 vs 0.823 unguided. The honest reading: (i) the damning comparison is local-vs-far, and “split the saturated cell into tighter sub-levels” is only worth its cost if the children actually move; (ii) guidance beats unguided displacement, but by little *here*, because this corpus occupies only a thin span (6 of 64 dimensions at the 95% energy level) and a random direction is already 82% transverse to it — guidance matters most when the corpus has already spread, which is exactly the late-run regime where refinement fires.

## II.7 The honest denominator

Recursive refinement creates an accounting problem that a coverage paper must not paper over: *refinement changes the measure being covered*. The refine policy's single run admits three defensible coverage numbers (mid radius, redteam / persona):

| denominator                                                                 | covered fraction |
|-----------------------------------------------------------------------------|------------------|
| base reachable pool (what the generator could do before refinement)         | 0.492 / 0.501    |
| the policy's own final reachable pool (base + everything refinement opened) | 0.960 / 0.971    |
| the full latent space (oracle: all modes, reachable or not)                 | 0.381 / 0.538    |

Each number is true and each, alone, misleads. Against the *fixed base pool*, refinement looks like a small regression (0.49 vs 0.54 for gated coverage without refinement): the budget it spent opening new territory was budget not spent filling old territory. Against the *full latent space* it looks like the only policy that works (0.381 vs 0.214; 0.538 vs 0.272) — but no base policy could have scored here at any budget, since 35 of 60 (20 of 40) modes were unreachable before refinement minted specs for them. And against its *own final reachable pool* it looks nearly finished (0.96+), which is partly real (tight child specs are easy to cover) and partly the denominator being shaped by the same process being graded — the refined proposal distribution concentrates exactly where the policy then samples.

The operational face of the problem is the pool refresh of §4: when refinement opens a region, the estimation pool must be re-drawn to include it, or the selector sees zero marginal gain precisely where the new territory is. We hit this live before we hit it in theory: in the pilot's first run, axis-conditioned candidates embedded  $\approx 0.88$  from an unconditioned-only pool with  $\epsilon_{\text{sel}} = 0.73$  — every candidate scored zero and the selector went blind — and the fix was to build the reference pool from *both* strata of the proposal process (§8). The denominator must track the reachable distribution of the pipeline as it currently exists.

Our reporting rule, followed in every table above: state the denominator next to every coverage number, and report at least (i) the fixed pre-refinement pool, which is comparable across policies and cannot be gamed by growing the space, and (ii) the policy's own final reachable pool, labeled as self-referential, with the growth of the reachable set (modes reached, spec count) reported alongside. A single “we covered X%” from a self-expanding pipeline reflects a choice of denominator.

## II.8 Live pilot

**II.8.1 Setup and what it cost** The pilot builds a red-team *evaluation* prompt suite on the real stack: gpt-5.6-luna via OpenRouter for generation, judging, attractor mining and axis elicitation; local nomic-embed-text (768-d, unit-normalized) for embeddings. Severity is deliberately mild — the items are benign user messages that probe whether an assistant adds disclaimers, resists sycophancy and leading questions, and

stays calibrated — and the judge scores a benign dimension that hard-gates acceptance. All state is append-only JSONL under `live_logs/` and the run is resumable; it was in fact resumed mid-flight after a directory move, which is what the checkpointing is for.

The run reached its full target: **50/50 accepted items** in 50 steps ( $K = 3$  candidates each), against a **50-item naive baseline** and a **360-item reference pool**. The pool has two strata — 240 unconditioned one-shot draws and 120 axis-conditioned but *unselected* draws — for a reason discovered the hard way (§8.3). Six axes were elicited, two refinements fired (8 axes final), and two ledger rounds mined attractors that became negative constraints. Judge quality averaged 8.89/10. Cost of the final instrumented segment: **74 calls, 36,056 prompt + 74,681 completion tokens, \$0.0968**, 394 s wall clock. Earlier segments (pool construction, the naive baseline, and the first 32 accepts) predate the usage ledger and are not included in that figure; total pilot spend was well under the \$3 budget but only the final segment is *measured*, so we report that number rather than an estimate.

**II.8.3 The blind-selector failure (why the pool has two strata)** The first version of this pilot used a reference pool of unconditioned draws only. Every axis-conditioned candidate then scored a marginal gain of *exactly zero*: measured post-hoc, method items sat at median distance 0.88 from the unconditioned pool while  $\epsilon_{\text{sel}}$  was 0.73, so no candidate’s ball contained any pool point and the selector was choosing uniformly at random among ties for 30+ consecutive steps. The bug is conceptual, not numerical: if the reference measure does not include the region the proposal distribution actually reaches, coverage-greedy selection is blind exactly where the method is working. The fix — pool the unconditioned and axis-conditioned strata — is the live counterpart of the simulation’s pool-refresh-after-refinement (§4), and it is the single most transferable lesson of the pilot: *the reference pool must be drawn from the same proposal process the policy uses, including every conditioning mechanism*.

**II.8.4 Literal vs latent diversity, and a kernel caveat** Tracking both n-gram and embedding diversity as the corpus grows reproduces a divergence reported at much larger scale on this model family: the two move in **opposite directions**.

| n  | distinct-1 | distinct-2 | Vendi (centered) | Vendi (uncentered) |
|----|------------|------------|------------------|--------------------|
| 5  | 0.614      | 0.944      | 3.98             | 3.62               |
| 10 | 0.482      | 0.861      | 8.01             | 5.04               |
| 20 | 0.367      | 0.809      | 16.16            | 7.37               |
| 30 | 0.338      | 0.791      | 23.65            | 10.00              |
| 40 | 0.312      | 0.772      | 30.55            | 11.87              |
| 50 | 0.295      | 0.770      | 37.06            | 13.54              |

Literal diversity falls monotonically (distinct-2 0.944  $\square$  0.770) while latent diversity rises monotonically (centered Vendi 3.98  $\square$  37.06). Reporting either alone supports an opposite conclusion about the same corpus. The mechanism is plain enough — a growing corpus reuses vocabulary and sentence frames while still spreading in meaning-space — but it means “diversity” must be qualified by level, and a coverage objective defined on embeddings is silent about literal repetition. (Against the naive baseline the method is slightly *less* literally diverse at distinct-1, 0.295 vs 0.341, and slightly more diverse latently, centered Vendi 37.06 vs 30.34, with markedly lower self-repetition, max 3-gram Jaccard 0.036 vs 0.083.)

**The uncentered Vendi is a kernel artifact.** Same-domain embeddings share a mean direction, and an uncentered linear kernel measures that shared cone as if it were content. On this data the two readings differ by 2.7 $\times$  (13.54 vs 37.06) and, worse, they order the growth curve differently in slope. We report both everywhere and treat only *relative* comparisons at fixed  $n$  as meaningful for the uncentered form.

**Is  $\epsilon$  measuring the cone or the content?** This is the question the caveat raises for every coverage number computed with cosine distances, so we measured it. Our pool’s mean pairwise cosine is **0.444** (p05 0.330, p95 0.591) — a wide cone, not the  $\sim 0.9$  concentration seen in some same-domain embedding sets — with

mean direction norm 0.667. Our  $\epsilon_{\text{sel}} = 0.770$  corresponds to cosine similarity 0.704, which sits at the **1.8th percentile** of pairwise distances: an  $\epsilon$ -ball at this radius captures genuine near-duplicates, not the bulk of the cone. So for this pilot the coverage numbers are measuring content. We record the diagnostic because the answer is embedder- and domain-dependent: with a tighter cone the same calibration rule would have produced an  $\epsilon$  inside the cone's width, and the resulting "coverage" would have been a restatement of the embedding geometry. Any coverage result computed on cosine distances should publish this check.

## II.9 Related work

**Submodular maximization.** The  $(1 - 1/e)$  greedy guarantee for monotone submodular objectives is Nemhauser, Wolsey & Fisher (1978); hardness of beating it for max-coverage is Feige (1998). Lazy greedy (Minoux 1978) and stochastic greedy (Mirzasoleiman et al. 2015) make greedy cheap; sieve-streaming (Badanidiyuru et al. 2014) gives the single-pass regime our sequential setting approximates. Facility location — of which our pool objective is the thresholded special case — is the standard submodular model of "representativeness" in data summarization (Lin & Bilmes 2011).

**Packing vs covering, k-center vs k-median.** Farthest-point traversal (Gonzalez 1985) 2-approximates k-center — the packing objective the sibling paper streams — while k-median/facility-location minimizes *average* service distance and behaves like our coverage objective (measure-weighted, bulk-first, outlier-indifferent). Coreset constructions (Har-Peled & Mazumdar 2004) make the same distinction: an  $\epsilon$ -coreset for k-median weights by mass, a k-center net does not. Our contribution is not new theory for these objects but the identification of *which one the corpus-construction product requirement actually is*, plus measurement of what optimizing the wrong one costs.

**Quality-diversity search.** MAP-Elites (Mouret & Clune 2015) maintains an "illuminated map": one elite per cell of a hand-designed behavior space — equal allocation over a fixed grid, in our vocabulary, with quality as the within-cell objective. Our method differs in that cells are elicited from the generator and refined recursively rather than fixed, allocation is marginal-gain-driven rather than one-per-cell, and coverage is measured against a reachable distribution rather than a designer-set grid; §6.5's equal-allocation row quantifies what the fixed-grid choice costs at small  $\epsilon$ .

**Diversity metrics and dedup.** The Vendi score (Friedman & Dieng 2022) is our spectrum-level diversity readout, computed at  $O(nD^2 + D^3)$  via the Gram trick as in Part I. SemDeDup (Abbas et al. 2023) removes semantic near-duplicates from web-scale corpora — the deletion-side mirror of our selection problem: both spend a budget to maximize non-redundant semantic territory. DPPs (Kulesza & Taskar 2012) give probabilistic diverse selection; exact DPP sampling scales poorly in  $n$  and optimizes a determinant (volume) rather than covered measure, though greedy MAP-DPP behaves similarly to facility-location greedy in practice.

**Eval-suite construction.** Perez et al. (2022) generate evaluations with LMs at scale; red-teaming work (Ganguli et al. 2022) documents exactly the clustering-into-attack-families failure that motivates a coverage objective for safety suites.

**Synthetic instruction corpora, and the baselines we compare against.** Self-Instruct (Wang et al. 2023) bootstraps instructions from a seed pool with few-shot exemplars and a ROUGE-L similarity filter; the released Stanford Alpaca corpus (Taori et al. 2023, 52,002 items) is its canonical output. Evol-Instruct / WizardLM (Xu et al. 2023) instead *mutates* existing instructions with depth and breadth evolution operators, released as WizardLM\_evol\_instruct\_V2 (143,000 items). Persona-Hub (Ge et al. 2024) conditions generation on a catalogue of  $\sim 1\text{B}$  personas, releasing 50,000 persona-synthesized instructions; it is the closest published relative of latent-axis conditioning, differing in that its catalogue is mined at scale and flat, where our axes are elicited from the generator and refined into a tree. AttrPrompt (Yu et al. 2023) similarly conditions on attribute dimensions. We compare against the released artifacts of the first three (§11) rather than against reimplementations, since a reimplementation can be weak in ways that flatter us. On the selection side our baselines are the standard ones for diverse subset choice: SemDeDup (Abbas et al. 2023), farthest-point traversal (Gonzalez 1985), and greedy MAP inference for DPPs (Kulesza & Taskar 2012; Chen, Zhang & Zhou 2018).

## II.10 Relation to Part I

The two papers are a packing / covering dual pair, and the duality is the k-center / k-median one:

|                   | infinite horizon (sibling)                 | finite budget (this paper)                  |
|-------------------|--------------------------------------------|---------------------------------------------|
| horizon           | unbounded stream                           | known $n = 10,000$                          |
| objective         | max-min gap + typicality anchor            | $\epsilon$ -ball covered measure of $\mu$   |
| classical problem | k-center / packing                         | max coverage / facility location            |
| optimum's habitat | boundary, extremes                         | bulk, measure-weighted                      |
| junk incentive    | attracted (needs gate badly)               | indifferent (gate still needed for quality) |
| planning          | impossible (no $n$ )                       | allocation across cells (§3.6)              |
| guarantee         | none exact (packing is 2-approx streaming) | $(1 - 1/e)$ , $\frac{1}{2}$ streaming       |
| refinement's role | capacity so the stream never exhausts      | reach — and a denominator liability         |

The method stack transfers almost unchanged — elicited axes, judge gate, ledger, recursive refinement are all objective-agnostic scaffolding for steering a generator; only the selection rule and its state differ (running second-moment matrix and min-gap scales there; reference pool and covered mask here). The sibling's slice-world results are directly load-bearing for this side of the duality. At an identical budget of 10,000 generations, its measured policy ladder runs from fixed-prompt (Vendi 1.06, reachable dimension 13) through random-spec (6.70) and score-guided specs (7.93) to axis scoring plus transversality-guided refinement (Vendi 11.23, reachable dimension 24) — a  $10.6\times$  diversity gain from conditioning and refinement alone, with refinement raising the reachable dimension from 13 to 24 (transversality theorem P4 in action). Two of its negative results transfer with particular force here. First, *naïve refinement hurt*: children split isotropically near their saturated parents concentrated probability mass exactly where the corpus already sat; refinement pays only when the new directions are transverse to the occupied span — our §6.8 reproduces this under the coverage objective (density, not reach). Second, its *inability-to-be-novel* metric was best ( $-0.104$ ) for the fixed-prompt policy precisely because that policy covered almost nothing and so consumed no headroom. The coverage-side analogue of that trap: a method that concentrates every  $\epsilon$ -ball in one small region leaves the rest of the space “available”, which flatters any residual- or headroom-based metric. This is why every table in §5 reports covered-fraction *jointly* with a within-corpus diversity measure (Vendi) and why coverage is always measured against the reachable manifold's measure, never against what happens to be left over. Its junk-conflation measurement is also starker than ours and worth quoting for eval-suite builders: in a hard-floor novelty bank with no typicality gate, the junk fraction of accepted items rose from 7.9% in the first half of the run to 48.2% in the second — once legitimate space saturates, the only candidates still clearing a novelty constraint are the incoherent ones. A red-team suite or persona corpus built that way looks maximally diverse by every embedding metric and is half garbage, silently.

The deepest difference is epistemic: the infinite-horizon paper never has to say what fraction of anything it achieved, while a coverage claim is *always* a ratio, which is why the denominator problem (§7) exists only on this side of the duality. Conversely, the budgeted side gets something the unbounded side cannot have: a completion semantics. A coverage run can report “we are 62% done at resolution  $\epsilon$  and the marginal item now buys 0.1%” — a spend / stop signal no packing objective provides.

## II.11 Comparison with released instruction corpora

We compare against the released corpora of the three most-used synthetic-instruction methods — Alpaca (Self-Instruct, 52k), PersonaHub (50k), and WizardLM Evol-Instruct (143k) — on scale-free coverage of human-written instructions, together with a budget-matched ablation of our own configurations.

**II.10.1 Protocol** Human-written reference: databricks-dolly-15k, split into two disjoint 3,000-item halves — a STEER half that our method may read as embeddings on the selection side, and an evaluation half that nothing in any pipeline ever reads, and at which no corpus (ours or released) was ever aimed. All scores below are on the evaluation half. Estimator: Naeem et al. (2020) coverage / density with reference-side k-NN radii, reported as AUC over  $k \in \{3, 5, 10, 20\}$ ; radii are a property of the reference alone, identical for

every corpus, with nothing tunable per corpus. Every corpus is evaluated as a uniform random sample at matched  $n = 450$ . Our arms and the reimplemented baselines all receive the same 175 human seed tasks (the `seed_tasks.jsonl` Alpaca was built from), the same generator, and the same budget of 2,400 generator calls — rejection and selection losses are counted, not hidden. The generator never sees a reference word in any arm.

**II.10.2 Method** The submitted configuration (v4) is coverage-retrieval conditioning with every render kept:

1. **Aim by retrieval.** Sample an uncovered reference region (density-weighted,  $\propto 1/r^2$ ) and retrieve the nearest texts we already own — seed tasks plus our own corpus — as the few-shot exemplars. The reference enters as embeddings on the steering side only; retrieval over owned text substitutes for the missing inverse oracle.
2. **Radius-adaptive mode.** The target's own k-NN radius selects the prompt mode: tight radius (dense region)  $\rightarrow$  imitate the local task family closely; wide radius (sparse region)  $\rightarrow$  full diversification, with the nearest own outputs shown as explicit in-prompt negatives.
3. **Literal-channel spread.** Level-usage balancing across the elicited axis lattice (the headroom term applied at conditioning time) and a terse-register mandate.
4. **Keep everything.** Coverage is monotone in items; at a generation-matched budget, every discard is a permanent loss.

**II.10.3 Results** All twelve corpora, evaluation half, matched  $n = 450$ :

| rank | corpus                                        | AUC           | precision    |
|------|-----------------------------------------------|---------------|--------------|
| 1    | <b>ours (coverage-retrieval conditioning)</b> | <b>0.4441</b> | <b>0.973</b> |
| 2    | Alpaca (Self-Instruct, 52k)                   | 0.3722        | 0.87         |
| 3    | ours: selective (1-of-8)                      | 0.2591        | 0.93         |
| 4    | PersonaHub (50k)                              | 0.2532        | 0.78         |
| 5    | WizardLM Evol-Instruct (143k)                 | 0.2329        | 0.77         |
| 6    | Self-Instruct (reimpl., same seeds/budget)    | 0.2188        |              |
| 7    | ours: conditioning keep-all                   | 0.2127        |              |
| 8    | few-shot from seeds                           | 0.2114        |              |
| 9    | ours: orthogonalized conditioning             | 0.2102        |              |
| 10   | axis-conditioned, unseeded                    | 0.1008        |              |
| 11   | naive                                         | 0.1003        |              |
| 12   | high temperature                              | 0.0947        |              |

Our method places first, 19% above Alpaca, with the highest precision in the field, at roughly one-twentieth of Alpaca's generation budget. An interim measurement at 22% of budget ( $n = 528$ ) already scored 0.4276 against Alpaca's 0.3666 on the same protocol, so the result does not depend on final corpus size. PersonaHub and WizardLM score below our selective arm despite 20–60 $\times$  the scale.

Budget-matched ablation (full corpora, 2,400 generator calls each, same evaluation half):

| configuration               | kept  | AUC           | density | precision    | recall       |
|-----------------------------|-------|---------------|---------|--------------|--------------|
| <b>full method (v4)</b>     | 2,398 | <b>0.6597</b> | 1.637   | <b>0.960</b> | <b>0.178</b> |
| Self-Instruct (reimpl.)     | 1,833 | 0.3262        | 1.050   | 0.949        | 0.161        |
| few-shot from seeds         | 2,399 | 0.3246        | 1.052   | 0.935        | 0.152        |
| conditioning, keep-all      | 1,310 | 0.3147        | 0.704   | 0.864        | 0.068        |
| orthogonalized conditioning | 1,152 | 0.2941        | 0.718   | 0.873        | 0.101        |
| selective (1-of-8)          | 304   | 0.2591        | 1.009   | 0.928        | 0.099        |

The ablation attributes the margin: retrieval-aiming and the radius-adaptive mode carry it (rows 1 vs 2–4), few-shot anchoring to the human seeds accounts for most of the baselines' scores (row 3 vs the unseeded arms at 0.09–0.10), Self-Instruct's ROUGE filter adds 0.0016 over bare few-shot, and both selection (row 6) and orthogonalized conditioning (row 5) reduce coverage relative to keeping everything — consistent with §III of the companion analysis: coverage is monotone and measure-seeking, so discarding and occupied-span avoidance are each counter-productive under this objective, whereas both are load-bearing under max-min.

**II.10.4 Notes Reference-sample asymmetry.** Our method consumes a sample of the target distribution, as embeddings, on the steering side; the released corpora had no such input. This is the method's designed capability — coverage is always coverage *of* something, and a method that may specify the something should — but the like-for-like no-reference comparison is rows 2–5 of the ablation, where our no-STEER configurations sit at parity with the Self-Instruct family. The precise claim: given a specification of the space to cover, even one the generator never reads, retrieval-aimed density-adaptive conditioning covers it substantially better than the strongest seeded baseline covers it at twenty times the budget.

**Coverage and internal diversity are different objectives.** The winning corpus has the lowest mean-centered Vendi of its cohort (62.1; density 1.69): it spends items where the reference measure is, near-duplicates included, exactly as facility location prescribes. Ranking these corpora by a single internal-diversity score would rank the coverage winner last.

All arm logs, the seed file, both reference halves, and the evaluation harness are in the repository; every number carries a provenance tag.

## II.12 Limitations

- **Simulation realism.** Gaussian mixture worlds with a spec layer capture mode structure, skewed mass, junk, and limited base reach, but not the ways real embedding geometry misrepresents semantic distinctness ( anisotropy, hubness), nor judges whose errors correlate with novelty — the live pilot mitigates but at  $n = 50$ , not 10,000.
- **$\epsilon$  is chosen, not learned.** All guarantees are per- $\epsilon$ ; our calibration (quantiles of reachable NN distance) is a heuristic, and §6.5 shows the cost of optimizing at the wrong radius. A multi-resolution objective (integrating coverage over an  $\epsilon$  prior) is the obvious next step.
- **The pool is the measure.** Everything is relative to the reachable pool; behaviors the base generator cannot emit are invisible until refinement opens them (§7), and pool refresh is only as good as the refinement trigger.
- **Selection bias on the estimation pool.** We hold out evaluation pools in simulation, but the live pilot's selection and reporting share one 240-item pool (with the bias direction stated); a larger live run should split them.
- **Live scale.**  $n = 50$ , one seed, one model; the live pilot is evidence of mechanism, not of effect size at  $n = 10,000$ .
- **What the benchmark win is and is not.** Section 9.5 shows our *selection rule* beating five baselines on a shared candidate pool. That is a claim about selection, not about generation: the pool was built by other methods, and a selector cannot cover territory its pool never reached (§6.6). The reachability gap remains the binding constraint, and closing it is a generation problem.
- **Single embedder, single similarity geometry.** Every coverage number here is computed in nomic-embed-text space with cosine distance. The parent project measured a correlation of only  $r = 0.155$  between text-embedding and CLIP-image-embedding pairwise similarity on rendered DALL-E outputs, which means text-space coverage is close to blind to whether the *rendered artifacts* are diverse. For any domain with a downstream rendering step, the objective should be defined in the space the artifact actually lives in; we did not do that here.
- **Matched-n does not neutralize generator quality.** In the instruction head-to-head (§11) the released corpora are 20-70x larger than ours and were produced by different, mostly stronger generators, with human curation in at least one case. Matched-n sampling controls for size; it cannot control for the

model that wrote the items, and PersonaHub's persona catalogue is orders of magnitude larger than any axis lattice we elicit.

## II.13 Conclusion

At a finite budget, corpus construction is a covering problem, and treating it as one pays: the coverage functional is exactly submodular in its Monte-Carlo form, greedy selection is near-optimal in theory and measurement, and the resulting suites blanket the reachable space instead of decorating its boundary. Packing — the right shape for an unbounded stream — is measurably the wrong one here, sacrificing 2–4× covered fraction to win a min-gap metric no budgeted product requirement asks for. The sibling stack's scaffolding (elicited axes, gates, ledger, recursive refinement) transfers whole; what does not transfer is innocence about denominators: a method that grows the space it is covering must say so when it reports how much of that space it covered.

---

py, the accompanying code, and the accompanying code in this directory and recorded in figures/summary\_sim.json, figures/summary\_theory.json, and figures/summary\_live.json. Reproduction: README.md'.

---

## Part III — Where the two objectives disagree

Everything above was two papers sharing a theory section. This part is the reason to publish them together: run both objectives on the same generator, the same embedder and the same budget, and they do not merely differ in emphasis — they rank methods almost oppositely.

### III.1 The dissociation, measured

In a leakage-free selection benchmark (one shared candidate pool, budget matched at 700 and 1,896 generations respectively, selector shown only an estimation half of the reference and scored on a disjoint held-out half):

| selector                       | DALL·E coverage | psychometric coverage | DALL·E min-gap |
|--------------------------------|-----------------|-----------------------|----------------|
| coverage-greedy (ours)         | <b>0.4829</b>   | <b>0.3901</b>         | 0.299          |
| coverage-stream (ours)         | 0.4200          | 0.3785                | 0.283          |
| ROUGE-L filter (Self-Instruct) | 0.4629          | 0.3575                | 0.208          |
| SemDeDup                       | 0.4600          | 0.3575                | 0.336          |
| random                         | 0.4629          | 0.3533                | 0.198          |
| MAP-DPP greedy                 | 0.3829          | 0.0988                | 0.402          |
| k-center (Gonzalez)            | 0.4114          | 0.0557                | <b>0.491</b>   |

**k-center wins min-gap outright in both domains and finishes last on coverage**, at 0.0557 against random's 0.3533 in the psychometric domain — a sixth of what random selection achieves. The two highest-Vendi selectors (k-center and MAP-DPP) are the two worst covering ones. A practitioner who reads "diversity" off a Vendi score and deploys the selector that maximizes it will get a corpus that covers less of the space than picking at random.

### III.2 Why the scoring rule has to fork

The four factors are shared but two of them invert:

|                    | max-min / packing                             | coverage / covering                                 |
|--------------------|-----------------------------------------------|-----------------------------------------------------|
| the ask            | no two items alike, in $n$ turns              | reach as much as possible, in $n$ turns             |
| governed by        | the closest pair                              | the bulk                                            |
| submodular         | no                                            | yes, so greedy carries $(1 - 1/e)$                  |
| classical kin      | k-center                                      | facility location                                   |
| <b>spread(a)</b>   | <b>min</b> distance between level centroids   | <b>measure-weighted mean</b> distance               |
| <b>headroom(a)</b> | levels not yet <b>used</b>                    | measure not yet <b>covered</b>                      |
| value selection    | <code>farthest_levels</code> (max-min subset) | <code>coverage_levels</code> (greedy marginal gain) |
| failure mode       | chases outliers into off-manifold junk        | leaves the frontier empty                           |

The behavioural difference is visible on a two-line test. Given an axis with three tightly-clustered common levels and one rare outlying level, `farthest_levels` picks the outlier **first** and `coverage_levels` picks a cluster centre first and the outlier second. That is k-center versus facility location, reproduced inside the axis axis scoring, and it is the mechanism behind the table above.

### III.2b The duality, taken seriously

Calling coverage “the dual of max-min” is exactly right in one sense and importantly wrong in another, and both halves have algorithmic consequences we verified on real embeddings.

**Where they are dual.** For radius  $\epsilon$ , a *maximal*  $\epsilon$ -packing is automatically an  $\epsilon$ -covering: any uncovered point could have been added to the packing. We confirm this numerically on 3,000 held-out human instructions — at every  $\epsilon$  tested, the greedy maximal packing covers 100.0% of the set, and the classical sandwich  $N_{\text{cov}}(\epsilon) \leq N_{\text{pack}}(\epsilon) \leq N_{\text{cov}}(\epsilon/2)$  holds. Max-min run to maximality is a covering algorithm, and greedy k-center 2-approximates the covering **radius**.

**Where they are not.** That duality concerns the worst-case radius: the distance from the least-covered point to its nearest center. The coverage that matters for synthetic data is **measure-weighted** — what fraction of the reference distribution lies within reach — and the two come apart exactly when the measure is non-uniform, which it always is. k-center is driven by outliers: every isolated point sets the max and so commands a center. Measure-weighted coverage is driven by mass: an outlier is worth its own measure and no more. At a budget far below what the space needs, the prescriptions are opposite — and that is the entire content of our k-center result (first on min-gap, last on coverage). The true dual of measure-weighted coverage is fractional set cover, whose primal is submodular; max-min is not.

**What the duality still buys: packing calibrates, covering places.** Even where packing is the wrong way to *place* points, it is the right way to *choose the scale*. Define  $\epsilon^*(n)$  as the radius at which a maximal packing of the reference has about  $n$  points — the finest resolution an  $n$ -item corpus can deliver. Measured on the held-out human reference:  $\epsilon^*(100) = 0.549$ ,  $\epsilon^*(450) = 0.472$ ,  $\epsilon^*(1000) = 0.418$ . This retroactively diagnoses our failed fixed- $\epsilon$  comparison with more precision than “the  $\epsilon$  was too small”: its  $\epsilon$  sat at the 2nd percentile of pairwise distances *regardless of budget*, while a 450-item corpus genuinely supports about the 2.6th percentile — the error was choosing resolution independently of  $n$ . Protocol: calibrate  $\epsilon$  by packing, then place by greedy marginal coverage against the measure.

**The ceiling decomposition.** In a generative setting, selection cannot cover what generation never proposed. We therefore report every coverage number as achieved = ceiling  $\times$  efficiency, where the ceiling is the coverage of the reference by the union of *all* candidates the generator produced, and efficiency is the fraction of that ceiling the selector captured. A low ceiling with high efficiency says the axes are too narrow and more search is wasted — only support expansion (recursive refinement) can help; the reverse says the selector is the problem. Without the split, “we lost on coverage” is not actionable.

### III.2c The comparison that proves the fork

Part II’s §11 compares our coverage method with the released Alpaca, PersonaHub and WizardLM corpora under a controlled protocol (same 175 human seeds, same generator, budget matched in generator calls,

evaluation on a reference half nothing ever read). Ours places first of twelve — 0.4441 against Alpaca's 0.3722, at one-twentieth of Alpaca's budget, with the field's highest precision. Two rows of its ablation matter to this joint paper beyond the ranking.

First, **the orthogonalization row**. The configuration applying Part I's orthogonalized conditioning to the coverage objective scores 0.2941, below plain conditioning's 0.3147: orthogonalization steers generation away from the occupied span, which is away from the reference-dense core coverage is paid to fill. Each part's load-bearing tool, applied to the other's objective, reduces performance. The two scoring rules fork because they must.

Second, **the Vendi row**. The winning coverage corpus has the *lowest* mean-centered Vendi of any configuration in its cohort (62.1; density 1.69): it spends items where the reference measure is, near-duplicates included. A single "diversity score" would rank the coverage winner last. There is no one number; there are two objectives.

### III.3 What both objectives share

Both are limited by the same thing, and neither optimizer can fix it. Coverage of a reachable manifold and packing within one are both bounded by the reachable manifold's dimension, which is set by how far the prompt can move the generator — Part I §3. Both are vulnerable to the same failure, in which a novelty-seeking score is satisfied by off-manifold output, and both need the same typicality constraint to be well-posed. And both are measured through an embedder whose geometry can invert the result: we report a cross-corpus coverage comparison in which our worst corpus by every other measure — 19.8% exact duplicates — scores the **highest** coverage, 50× a published corpus's, because at an  $\epsilon$  in the 2nd percentile of reference distances the metric rewards centrality rather than spread.

### III.4 Practical guidance

1. **Say which objective you mean.** They are different problems with different optimal policies, and the vocabulary ("diverse", "varied") does not distinguish them.
2. **Count exact duplicates before computing anything.** A 72% duplicate rate makes every distance statistic a statistic about duplication.
3. **Publish the pairwise-similarity distribution** your kernel operates on. Mean pairwise cosine was 0.883 in one of our corpora and 0.444 in another with the same embedder; an  $\epsilon$  or a Vendi score means nothing without it.
4. **Measure at the level of the artifact you ship.** Text-embedding similarity explains ~2% of the variance in whether rendered images look alike.
5. **Report the procedure with the number.** Our own framework changed three times during this work; every result in this paper is tagged with the procedure that produced it.